

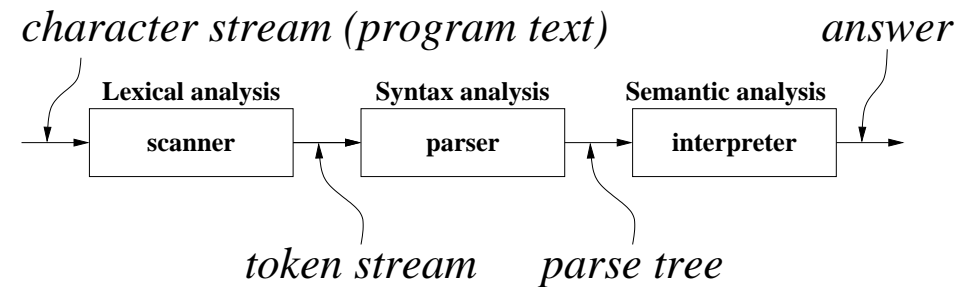
Copyright (C) 2021 Timothy Fossum

Permission is granted to copy, distribute and/or modify this document under the terms of the GNU Free Documentation License, Version 1.1 or any later version published by the Free Software Foundation. A copy of the license is included in the file "COPYING", entitled "GNU Free Documentation License".

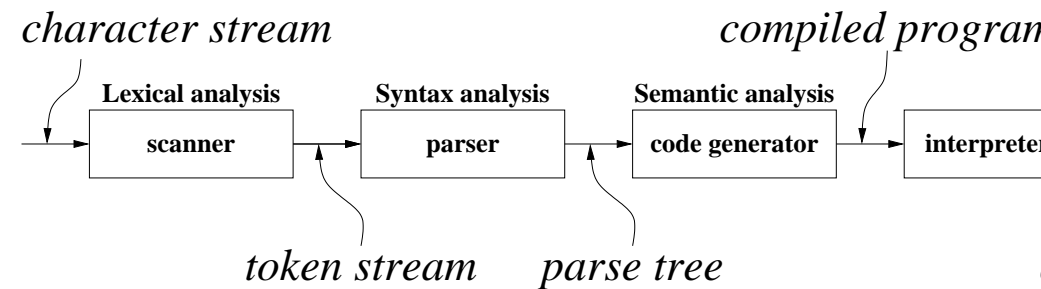
## Environment-Passing Interpreters

Interpretation vs. compilation can be illustrated by a picture:

Interpreter execution:



Compiler execution:



## Environment-Passing Interpreters (continued)

Most programming languages have grammar rules defining an *expression*. In principle, an expression typically involves values (like variables, integers, and the like) and operators (like addition and multiplication). Example Java expression: `foo(11) && toggle`. In all of the languages we discuss in this class, every expression has a value. Such languages are called “expression-based languages”.

An *expressed value* is the value of an expression as specified by the language. For example, the expressed value of the Java expression `‘2+3’` is 5. A *denoted value* is a value denoted by a symbol. Denoted values are internal to the interpreter, whereas expressed values are values of expressions that can be seen “from the outside”.

For a symbol, say `x`, you normally think that the value of the expression `x` is the same as the value of `x`. But what about a language such as Java? In Java, the denoted value of a variable is a *reference* to an object, whereas the expressed value of the variable is the object itself. This may seem like a subtle distinction, but you will see its importance later.

In summary, for a symbol, its expressed value is what gets displayed when you call its `toString` representation, for example), and its denoted value is the value bound to the symbol in its environment. In our early languages, the denoted values and expressed values were the same. In our later languages, we will see why we need to separate denoted values from expressed values to implement language features such as mutation.

## Environment-Passing Interpreters (continued)

You should also distinguish between a *source language* and its *implementation language*. A source language is a language to be interpreted, and its implementation language is the language in which the interpreter is written. (The term *source language* is often used to refer to a source language. Similarly, the term *implementation language* is often used to refer to its implementation language.)

In the rest of this course, our source languages will be a collection of languages used to illustrate the various stages of language design, and our implementation language will be Java. Don’t be disappointed by the term “source language”; the languages we define have significant computational power, and they will illustrate a number of core ideas that are present in all programming languages.

We start with a language we call “Language V0” – think of this as “Version Zero”. Its grammar specification file appears on the next slide.

## Language V0

```
# Language V0
skip WHITESPACE '\s+'
LIT '\d+'
ADDOP '\+'
SUBOP '\-'
ADD1OP 'add1'
SUB1OP 'sub1'
LPAREN '\('
RPAREN '\)'
COMMA ','
VAR '[A-Za-z]\w*'
%
<program> ::= <exp>
# these three grammar rules define what it means to
<exp>:LitExp ::= <LIT>
<exp>:VarExp ::= <VAR>
<exp>:PrimAppExp ::= <prim> LPAREN <rands> RPAREN
<rands> ::= <exp> +COMMA
<prim>:AddPrim ::= ADDOP
<prim>:SubPrim ::= SUBOP
<prim>:Add1Prim ::= ADD1OP
<prim>:Sub1Prim ::= SUB1OP
%
#include code
```

## Language V0 (continued)

Observe that this PLCC language specification has all three of the in Slide Set 1a: the lexical specification section (token specification section (given as BNF rules), and the semantic specification (just an `#include` line).

You can consider a grammar rule like

```
<program> ::= <exp>
```

to mean that “a program consists of an exp” (where exp stands for expression). Similarly, you can consider a grammar rule like

```
<exp>:LitExp ::= <LIT>
```

to mean that “a LitExp is (an instance of) an expression (an exp) of a LIT.” Similar remarks apply to all the grammar rules.

The ‘`#include code`’ line at the end of this file means that the file named `code` should be inserted into the grammar file input point to be processed by PLCC.

## Language V0 (continued)

Example “programs” in this language:

```
3
x
+(3, x)
add1 ( +(3, x) )
+(4, -(5, 2))
```

Observe that in Language V0 – and in most of the other languages in these class notes – we write arithmetic expressions in *prefix form*, where the arithmetic operator (such as ‘+’ or ‘-’) precedes its operands. (A prefix form like ‘+(3, x)’ would normally be written mathematically as ‘3+x’ in *infix form*. It turns out that prefix form expressions are easier to parse than infix form expressions, which is why we use prefix form in our Language V0. See the INFIX language in Slide Set 6 for a further discussion of infix form.

Prefix form is not entirely unusual: languages in the Lisp family (including Scheme) use prefix form. Contrast this to languages such as C, Java, and Python, where arithmetic operators appear principally in infix form.

## Language V0 (continued)

Here is a mapping from the concrete (BNF) syntax of Language V0 to its abstract syntax tree representation as an abstract syntax. The Java class files are created automatically by the PLCC. Each item in a Box is the signature of the corresponding class.

<program>	::= <exp> <span style="border: 1px solid black; padding: 2px;">Program(Exp exp)</span>
<exp>:LitExp	::= <LIT> <span style="border: 1px solid black; padding: 2px;">LitExp(Token lit)</span>
<exp>:VarExp	::= <VAR> <span style="border: 1px solid black; padding: 2px;">VarExp(Token var)</span>
<exp>:PrimappExp	::= <prim> LPAREN <rands> RPAREN <span style="border: 1px solid black; padding: 2px;">PrimappExp(Prim prim, Rands rands)</span>
<rands>	**= <exp> +COMMA <span style="border: 1px solid black; padding: 2px;">Rands(List&lt;Exp&gt; expList)</span>
<prim>:AddPrim	::= ADDOP <span style="border: 1px solid black; padding: 2px;">AddPrim()</span>
<prim>:SubPrim	::= SUBOP <span style="border: 1px solid black; padding: 2px;">SubPrim()</span>
<prim>:Add1Prim	::= ADD1OP <span style="border: 1px solid black; padding: 2px;">Add1Prim()</span>
<prim>:Sub1Prim	::= SUB1OP <span style="border: 1px solid black; padding: 2px;">Sub1Prim()</span>

### Language V0 (continued)

The term *abstract syntax* might seem odd because it refers to a collection of explicit Java classes. Instead, the term *abstract* here means that we keep only the information on the right-hand-side (RHS) of the grammar rules that can change, principally by ignoring certain RHS tokens. For example, the `<exp>:PrimappExp` grammar rule has tokens `LPAREN` and `RPAREN` on its LHS and RHS, but the generated `PrimappExp` class does not have fields for these tokens: they are “abstracted away”.

Because the `<exp>` and `<prim>` nonterminals appear on the LHS of the grammar rules, we must disambiguate these grammar rules by annotating the nonterminals with appropriate class names. For these grammar rules, the nonterminal corresponds to the name of an abstract (base) Java class that is obtained by capitalizing the first letter of the nonterminal name. For example, the classes `LitExp`, `VarExp`, and `PrimappExp` extend the abstract base class `Exp`. Similarly, the `AddPrim`, `SubPrim`, `Add1Prim` and `Sub1Prim` classes extend the abstract base class `Prim`.

Once you create the grammar specification file and run `plccmk`, it generates the Java code in the `Java` subdirectory. Here you can see, for example, that the `LitExp` class extends the `Exp` class and that the `AddPrim` class extends the `Prim` class.

### Language V0 (continued)

The `Program` class has one instance variable named `exp` of type `Exp`. Since `Exp` is an abstract class, an object of type `Exp` must be an instance of a class that extends `Exp`: namely, an instance of `LitExp`, `VarExp`, or `PrimappExp`. The `Program` class does not have a constructor, so you can’t instantiate an object of type `Program`.

The directory `/usr/local/pub/plcc/Code/V0` contains the grammar specification file, named `grammar`, for this language.

The `grammar` file in Language V0 has three parts, separated by lines that begin with the single ‘%’: the **lexical specification section**, the **syntax specification section**, and the **code (semantics) section**.

Recall that if your grammar file has only the lexical specification section, the `plccmk` tool produces Java code for a scanner (`Scan`), but nothing else. If your grammar file has only the lexical specification and syntax specification sections, the `plccmk` tool produces Java code for a scanner (`Scan`) and a parser (`Parser`) for the grammar, but nothing else.

## Language V0 (continued)

The code section defines the language semantics. In this section, the behaviors defined by the grammar rules are given life by defining their behavior by defining the `$run()` method in the start symbol class (`Program` of language V0).

In the absence of a redefined `$run()` method – for example, if the method is omitted – the default `$run()` behavior in the `_Start` class is to print the value of the start symbol in a format illustrated in Slide 1.29.

In later versions of this language (V1 and beyond), we see how the `$eval` method can be used to print the arithmetic value of an expression. But in V0 we will be content with simply printing a copy of the expression.

*In the code section of a PLCC language specification, the behavior of the `$run()` method in the start symbol class defines the language semantics.*

## Language V0 (continued)

Assuming that we have created the grammar file in a directory named `grammar`, running the `plccmk` tool creates a Java subdirectory with source files named `Program.java`, `LitExp.java`, and so forth, that correspond to the syntax classes shown in Slide 3.7. In the Java directory, you can find the source files named `Token.java`, `Scan.java`, `Parse.java`, and `Rep.java`.

The `Rep` program repeatedly prompts you for input (with ‘-->’), parses the input, and prints the result: a `String` representation of the expression with all white space removed. If you want to run this program from the directory containing the grammar file – V0 in this case – you can run it as follows:

```
$ java -cp Java Rep
--> add1( + (2
    , 3))
add1(+ (2, 3))
--> ...
```

As we discussed in Chapter 1, parsing is the process by which a sequence of tokens (a *program*) can be determined to belong to the language defined by a grammar. We showed examples of leftmost derivations and how the derivations can be used to detect whether or not the program is syntactically correct.

### Language V0 (continued)

We get more than a success or failure response from our parser: *it returns a Java object that is an instance of the class determined by the grammar start symbol*. This object is the root of the *parse tree* of the program. It captures all of the elements of the parsed program.

In our Language V0 grammar, the start symbol is `<program>`, so the parse tree is an instance of the `Program` class.

### Language V0 (continued)

We have seen that the RHS of a grammar rule determines what instances belong to the class defined by its LHS. Only those entries on the RHS that appear in angle brackets `< . . . >` appear as instance variables; *any other RHS entries are token names that are used in the parse but that are abstracted away from the objects in the parse tree*.

For example, consider the following grammar rule in Language V0:

```
<exp>:PrimappExp ::= <prim> LPAREN <rands>
```

This rule says that `PrimappExp` is a class that extends the `Exp` class. The instance variables in this class are

```
Prim prim;  
Rands rands;
```

### Language V0 (continued)

Consider the following grammar rule in Language V0:

```
<exp>:LitExp ::= <LIT>
```

This rule creates a Java class `LitExp` having a single field name `Token`. The lexical specification defines a `LIT` token to be a sequence of one or more decimal digits.

Continuing in this way, each of the BNF grammar rules of Language V0 (see Figure 3.7) defines a class given by its LHS with a well-defined set of instances corresponding to the *<angle bracket>* entries in its RHS.

As we have already observed, when we encounter situations where two different nonterminals have the same (Token or nonterminal) name in angle brackets, we distinguish them by providing different instance variable names.

### Language V0 (continued)

Recall that repeating grammar rules have fields that are Java lists. For example, the Language V0 grammar has the following repeating rule:

```
<rands> **= <exp> +COMMA
```

This rule says that the `<rands>` nonterminal can derive zero or more expressions, separated by commas. The following sentences would match the `<rands>` nonterminal:

```
a, b, c      <-- <exp> COMMA <exp> COMMA <exp>
1, +(2, 3)   <-- <exp> COMMA <exp>
add1(x)      <-- <exp>
              <-- empty string
```

The class defined by this rule is named `Rands`. Its RHS shows only one nonterminal `<exp>`, so its corresponding Java class `Rands` has a field `exp` of type `List<Exp>`.

You might wonder how we chose a name like “rands”. It’s actually a shortened form of the term “operands”. In mathematics and in programming, we often refer to the things being operated on. For example, given the expression `+(2, 3)`, the operator is ‘+’ and its operands are 2 and 3. (Similarly, some languages use the term “rator” as a shortened form of the term “operator”.)



## Language V0 (continued)

When we parse a program such as

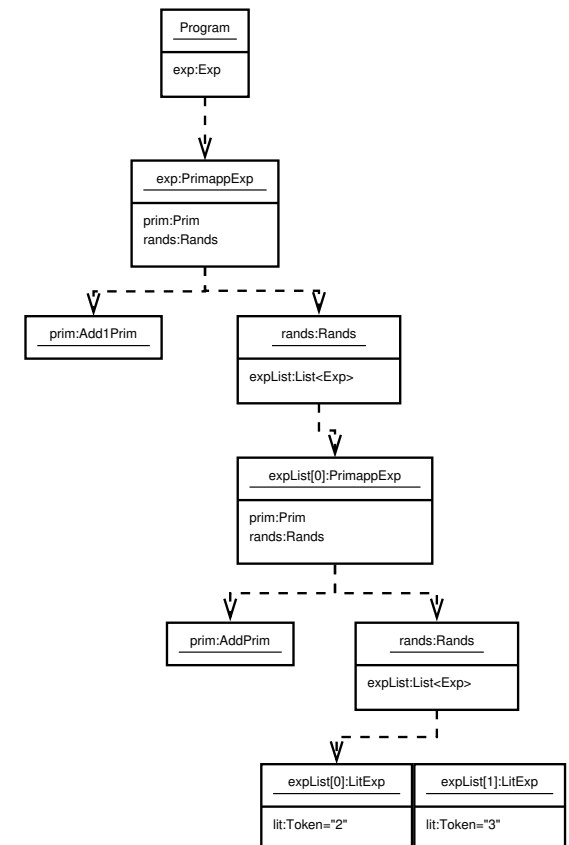
```
add1 ( + (2, 3) )
```

the parser returns an object of type `Program`. The `Program` object has an instance variable: `exp` of type `Exp`. The value of the `exp` instance variable is an object of type `PrimappExp` (which extends the `Exp` class) that has two instance variables: `prim` of type `Prim` and `rands` of type `Rands`. The value of the `prim` instance variable is an object of type `Add1Prim` (which extends the `Prim` class) that has two instance variables. And so forth ...

On the following slide we show the entire parse tree of this expression

## Language V0 (continued)

Parse tree for `add1 ( + (2, 3) )` in UML format:



## Language V0 (continued)

The Rep program defaults to running the `$run()` method in the `Program` class which prints the `Program` object as a `String`. (Actually, it prints the `Program` object as a `String`, but since the `Program` class extends the `Object` class, these behaviors are the same.) The default `$run` behavior – which if `$run` is not redefined in the `Program` class – displays a string like this:

```
Program@....
```

We want to show how to override the default `$run()` method in the `Program` class so that it prints the same text as the program input with extra spaces moved!

This means that we should see the following when interacting with the program from Language V0:

```
--> add1( +(2, 3) )
add1( +(2, 3) )
--> x
x
--> + ( p ,
- ( q, r) )
+ (p, -(q, r))
--> ...
```

## Language V0 (continued)

We follow the method described in Slide Set 1 to redefine the default `$run()` method in the `Program` class. *In all of our languages, the semantics of a program is the output produced by the `$run()` method on the root of the parse tree of the program.*

Recall that to add Java code (fields and method definitions) to a `Program`, use the following template:

```
Program
%%%
...Java code...
%%%
```

Since PLCC inserts this code verbatim into the body of the `Program` class, methods defined in this code can access all of the instance variables of the `Program` object. So for the `Program` object, these methods can refer to the `exp` instance variable of type `Exp`. Since the RHS of the `<program>` grammar rule is just the `$run()` method of the `Program` class is simple:

```
Program
%%%
    public void $run() {
        System.out.println(exp.toString());
    }
%%%
```

## Language V0 (continued)

To finish our implementation, we show how to implement the `toString()` method for an `Exp` object so that it returns a `String` representation.

There are three `Exp` classes: `LitExp`, `VarExp`, and `PrimappExp`. The `toString()` method for the first two classes is particularly easy, as they have right-hand sides that are just token strings: their `toString()` methods return the `String` value of the corresponding `Token` instance variable (see Figure 3.7):

```
LitExp
%%%
    public String toString() {
        return lit.toString();
    }
%%%

VarExp
%%%
    public String toString() {
        return var.toString();
    }
%%%
```

## Language V0 (continued)

Examine the rule for a `PrimappExp`:

```
<exp>:PrimappExp ::= <prim> LPAREN <rands> RPAREN
                        PrimappExp(Prim prim, Rands rands)
```

A `PrimappExp` object has just two instance variables:

```
Prim prim;
Rands rands;
```

There aren't any instance variables corresponding to `LPAREN` and `RPAREN`, because the `PrimappExp` class abstracts away these tokens. The only thing to do, then, is to re-insert them back into the `toString` result, in the same order as they appear on the RHS of the grammar rule. The `toString` method for `PrimappExp` calls `prim` and `rands` are called implicitly.

```
PrimappExp
%%%
    public String toString() {
        return prim + "(" + rands + ")";
    }
%%%
```

## Language V0 (continued)

Each of the <prim> rules has an RHS that corresponds to a Token by the parser. Just as we re-inserted the LPAREN and RPAREN tokens, we defined the toString method in the PrimappExp class, each of these classes simply returns the corresponding string token:

```
AddPrim
%%%
    public String toString() {
        return "+";
    }
%%%

SubPrim
%%%
    public String toString() {
        return "-";
    }
%%%

Add1Prim
%%%
    public String toString() {
        return "add1";
    }
%%%
```

The Sub1Prim code is similar and has been omitted.

## Language V0 (continued)

We have covered all of the plcc-generated classes except for Rands. We pay a bit more attention since a Rands object has an expList instance variable. It is a List of expressions. First examine the <rands> grammar rule:

```
<rands>  **= <exp> +COMMA
          Rands(List<Exp> expList)
```

To build a toString method for this class, we call the toString method on each of the expList entries and construct a String that puts them together. Here is the code:

```
Rands
%%%
    public String toString() {
        String s = ""; // the string to return
        String sep = ""; // no separator for the first expression
        // get all of the expressions in the operand list
        for (Exp exp : expList) {
            s += sep + exp; // exp.toString() is called
            sep = ","; // commas separate the remaining expressions
        }
        return s;
    }
%%%
```

## Language V0 (continued)

We can now re-build the Java code for this grammar using the `place` program. Assuming that everything compiles correctly, we should get the desired output from the `Rep` program: `Rep` parses each syntactically correct input and displays the resulting `Program` object as a `String`, which appears as the input with whitespace removed.

## Language V1

Now that you see how a parse tree for Language V0 can *print* itself, we see that a parse tree can *evaluate itself*.

The term *evaluate* can have many meanings (one of which is to produce a textual representation of itself), but for our purposes, to evaluate an arithmetic expression such as `add1( + (2, 3) )` means to produce the integer value of the expression. In other words, the value of an arithmetic expression is its numeric value using the rules for arithmetic.

(Remember that we are abstracting the notion of *value* to refer to an instance of the `Val` class. In this setting, a numeric value is an instance of the `Int` class of `Val`.)

If an expression involves an identifier (symbol), we need to determine the value bound to that identifier in order to evaluate the expression. For example, if the identifier `"x"` is bound to the integer value 10: then the expression `x + 9` would evaluate to 19.

*The interpreter evaluates every expression in some environment. The environment determines how to obtain the values bound to the identifiers that appear in the expression.*

## Language V1

The `Exp` class is the appropriate place to declare evaluation behavior using a method called `eval`. Here is how we declare the method in the (abstract) `Exp` class. Every class that extends `Exp` defines this method:

```
Exp
%%%
    public abstract Val eval(Env env);
%%%
```

Language V1 is the same as Language V0, except for adding `eval` classes that extend the `Exp` class. We continue to consider the only one to be an `IntVal` that holds an integer value.

Language V1 has three new files included in its grammar language: `envRN`, `val`, and `prim`.

- The `envRN` file contains Java class definitions to implement environments as discussed in Slide Set 2.
- The `val` file defines the `Val` class and its `IntVal` subclass – a simple wrapper for Java `ints`.
- The `prim` file defines the semantics of the seven `<prim>` BNF constructs in Language V1.

## Language V1 (continued)

Three classes extend the `Exp` class: they are `LitExp`, `PrimappExp`, and `AppExp`. We'll start with `LitExp`. Here is the code part of the `lit` file that defines the `eval` behavior of a `LitExp` object. (The `eval` method coexists with the `toString` behavior that we defined in Language V0; we do not show this here.)

```
LitExp
%%%
    public Val eval(Env env) {
        return new IntVal(lit.toString());
    }
%%%
```

Remember that a `LitExp` has a `Token` field named `lit`. When we call the `toString()` method on this field, we get the string of decimal digits from the part of the program text we are parsing. The `IntVal` constructor converts this into a real Java `int` that becomes part of the `IntVal` instance. The environment doesn't have anything to do with the value of a number literal; `10` evaluates to the integer value `10` no matter what environment we are in, so the `eval` routine for a `LitExp` simply returns the appropriate `IntVal`.

## Language V1 (continued)

Next we consider `VarExp`. Here is the code part of the grammar and the `eval` behavior of a `VarExp` object.

```
VarExp
%%%
    public Val eval(Env env) {
        return env.applyEnv(var);
    }
%%%
```

A `VarExp` object has a `var` instance variable of type `Token`. Given an environment, the value bound to `var` is precisely the value returned by `applyEnv` in turn is the value of the expression.

*The value of an expression consisting of a symbol is the value bound to the symbol in the environment in which the expression is evaluated, as determined by `applyEnv`.*

## Language V1 (continued)

Finally we consider `PrimappExp`. A `PrimappExp` object has two instance variables: a `Prim` object named `prim` and a `Rands` object named `rands`. To evaluate such an expression, we need to apply the given primitive operation (in the `Prim` object) to the values of the expressions in the `rands` object.

An object of type `Rands` has a `List<Exp>` instance variable named `expList`. In order to perform the operation determined by the `prim` object, we need to evaluate each of the expressions in `expList`. A utility method named `evalRands` in the `Rands` class does the work for us. Of course, this method needs to know what environment is being used to evaluate the expressions, so an `Env` is a parameter to this method.

```
Rands
%%%
    public List<Val> evalRands(Env env) {
        List<Val> args = new ArrayList<Val>(expList.size());
        for (Exp exp : expList)
            args.add(exp.eval(env));
        return args;
    }
%%%
```

## Language V1 (continued)

We *specify* that the expressions in an `evalRands` method call be evaluated to-last (or left-to-right, depending on how you are looking at it) or, not necessarily the case for all programming languages. In particular, see *et seq.* for a more complete discussion of order of evaluation.

The `evalRands` method returns a *list* of `Vals`. In order to access them easily and to apply normal arithmetic operations to them, we convert the *array* of `Val` objects. The utility method named `toArray` in the `Val` class accomplishes this.

*The expressions appearing in an application of a primitive are called also called **actual parameters**; the values of these expressions are called **arguments**.*

As a careful reader of these notes, you will have observed that the name `Rands` is derived from the word **operands**, and that the name `evalRands` method is derived from the word **arguments**.

## Language V1 (continued)

We now have the pieces necessary to define the `eval` method in the `PrimappExp` class:

```
PrimappExp
%%%
    public Val eval(Env env) {
        // evaluate the terms in the expression list
        // and apply the prim to the array of Vals
        List<Val> args = rands.evalRands(env);
        Val [] va = Val.toArray(args);
        return prim.apply(va);
    }
%%%
```

In summary, to evaluate a primitive application expression (a `PrimappExp` object), we first evaluate the operands (a `Rands` object) in the given environment and then pass the resulting argument array to the `apply` method of the primitive (a `Prim` object), which returns the appropriate value.

We are left to define the behavior of the `apply` methods in the `Prim` classes. Observe that by the time a `Prim` object gets its array of arguments, the environment no longer plays a role, since we have already evaluated all the expressions: only values remain.



## Language V1 (continued)

Since we are using the `apply` method with a `Prim` object, we need a declaration for this method to the (abstract) `Prim` class. Here is how we

```
Prim
%%%
    // apply the primitive to the passed values
    public abstract Val apply(Val [] va);
%%%
```

We use the parameter name ‘`va`’ to suggest the idea of a **value array**

A `Prim` object (there are seven instances of this class) has no instance methods. However, we can endow these objects with behavior, so that an `AddPrim` object knows how to add things, a `SubPrim` object knows how to subtract things, and so forth.

## Language V1 (continued)

Four of the `Prim` objects need two arguments (`+`, `-`, `*`, and `/`), and three of them need one argument (`add1`, `sub1`, and `zerop`). Since `Val` objects have `intVal()` arguments, we can grab the appropriate items from this array – the `args` array – and then, depending on the operation – to evaluate the result. Here is the `AddPrim` class:

```
AddPrim
%%%
    public Val apply(Val [] va) {
        if (va.length != 2)
            throw new PLCCEException("two arguments expected");
        int i0 = va[0].intVal().val;
        int i1 = va[1].intVal().val;
        return new IntVal(i0 + i1);
    }
%%%
```

The `intVal()` method calls shown in this code convert `Val` objects (e.g., `va[0]`) into `IntVal` objects – essentially like “downcasting”. The `IntVal` class, in turn, has Java `int` fields named `val`. So both `i0` and `i1` are `int` values. The `IntVal` class also has `+` and `-` methods that can be added together to return the resulting `IntVal` object. The `Prim` class defines the `intVal()` method behavior: an attempt to apply the `intVal()` method to a `Val` object that is *not* an `IntVal` throws an exception.

## Language V1 (continued)

The definitions of `apply` for the classes `SubPrim`, `MulPrim`, and `DivPrim` (the latter two are added in V1) have obvious implementations, except that the `apply` method throws an exception if it detects an attempt to apply a binary operator to a zero argument. This code is not shown here.

For the `Add1Prim` class, the `apply` method expects only one argument to be passed as element zero of the `va` array.

```
Add1Prim
%%%
    public Val apply(Val [] va) {
        if (va.length != 1)
            throw new PLCCEException("one argument expected");
        int i0 = va[0].intVal().val;
        return new IntVal(i0 + 1);
    }
%%%
```

Again, the definition of `apply` for the `Sub1Prim` class is entirely analogous. The definition of `apply` for the `ZeropPrim` class returns an `IntVal` of 0 for a zero argument and an `IntVal` of 1 (true) for a nonzero argument.

## Language V1 (continued)

In our final implementation step, we will define the `$run()` method of the `Program` object that displays the string representation of the *value* of its expression. We will show on the next slide for how we implement this.

An empty environment would only allow for integer expressions with no variables, since every variable would be unbound. To test Language V1, we create an environment `initEnv` specific to this language that has the following bindings (think Roman numerals!):

```
i => 1
v => 5
x => 10
l => 50
c => 100
d => 500
m => 1000
```

For Language V1, this environment can be obtained by a call to `Env.initEnv()`. In the `Program` class, we set the static `initEnv` variable to this environment. The bindings give us some variables to play with, though, as you will see when we dispense with them later.

## Language V1 (continued)

Here is the new `$run()` method for the `Program` object, along with the initial environment `initEnv` in the `Program` class:

```
Program
%%%
    public static Env initEnv = Env.initEnv();

    public void $run() {
        System.out.println(exp.eval(initEnv).toString());
    }
%%%
```

To test this, run the `Rep` program defined in the `Java` subdirectory, pressing `enter` at the prompts:

```
java -cp Java Rep
```

## Language V2

Language V2 is the same as Language V1, with the addition of semantics of an `if` expression. The relevant grammar rule and its representation are shown here:

```
<exp>:IfExp ::= IF <exp>testExp THEN <exp>trueExp ELSE <exp>>falseExp
IfExp(Exp testExp, Exp trueExp, Exp falseExp)
```

Notice that we need to add token names `IF`, `THEN`, and `ELSE` to our lexer, along with their obvious regular expressions.

The RHS of the `IfExp` grammar rule has three occurrences of the `<exp>` non-terminal. Since the `<exp>` items on the RHS of this grammar rule are non-terminals, we have named these instance variables of the class, we have named these instance variables `trueExp`, and `falseExp`, respectively. Each of these objects is an instance of the `Exp` class. The `IfExp` class has three instance variables:

```
Exp testExp;
Exp trueExp;
Exp falseExp;
```

**[Exercise (not to hand in):** See what would happen if you used '`<IF>`' instead of '`IF`'.]

## Language V2

To evaluate an `if` expression with a given environment, we first evaluate the `testExp` expression. If this evaluates to true, we evaluate the `trueExp` expression and return its result as the value of the entire expression. If the `testExp` evaluates to false, we evaluate the `falseExp` expression and return its result. In either case, the appropriate expression is evaluated in the given environment.

Since all instances of `Val` are really `IntVals` (for the time being), we define an `isTrue()` method for an `IntVal` object corresponding to 0 to be false and *all others* to be true.

## Language V2 (continued)

We define an `isTrue()` method for an `IntVal` object as follows. The following code is part of the `IntVal` class that is defined in the `val` file – only the `isTrue` method is given here:

```
public boolean isTrue() {  
    return val != 0; // nonzero is true, zero is false  
}
```

Observe that the `eval()` method in the `IfExp` class applies the `isTrue()` method to a `Val` object, so we must include a declaration for the `isTrue()` method in the `Val` base class. *Since we currently treat any `Val` object as true if it's not an `IntVal` of zero, our default `isTrue()` method in the `Val` class defaults to returning true.*

```
Val  
%%%  
...  
    public boolean isTrue() {  
        return true;  
    }  
...  
%%%
```

## Language V2 (continued)

```
<exp>:IfExp ::= IF <exp>testExp THEN <exp>trueExp ELSE  
                IfExp(Exp testExp, Exp trueExp, Exp falseExp)
```

Here is the eval code for the IfExp class:

```
IfExp  
%%  
    public Val eval(Env env) {  
        Val v = testExp.eval(env);  
        if (v.isTrue())  
            return trueExp.eval(env);  
        else  
            return falseExp.eval(env);  
    }  
%%
```

The `isTrue()` boolean method applies to any instance of `Val`. It is a method used only to implement the semantics of the `if...then...else` expression; it is *not* part of the source language. On the other hand, `zero?` is a *primitive* in the source language (starting with Language V1), not a method. The `zero?` primitive applies only to integer values in the source language. We define the *semantics* of the `zero?` primitive using the apply Java `ZeropPrim` class. You may find this somewhat confusing.

## Language V2 (continued)

```
<exp>:IfExp ::= IF <exp>testExp THEN <exp>trueExp ELSE  
                IfExp(Exp testExp, Exp trueExp, Exp falseExp)
```

Observe that the `eval` method in the `IfExp` class evaluates *only* the `trueExp` or `falseExp` expressions, never both. This is a semantic feature, not a syntax feature – of the definition of `eval` for an `if` expression. A *special form* refers to semantic structures that behave like expressions, but when evaluated, don't evaluate all of their constituent parts. An `if` expression is an example of a special form.

Some examples of `if` expressions are on the next slide.

## Language V2 (continued)

```
if 1 then 3 else 4
% => 3
```

```
if 0 then 3 else 4
% => 4
```

```
if
  if 1 then 0 else 11
then
  42
else
  15
% => 15
```

```
+(3, if -(x,x) then /(5,0) else 8)
% => 11 (note that the /(5,0) expression is not ev
```

You must understand that *an if expression is an expression and the evaluates to a value*. It is entirely unlike if statements in imperative languages like Java and C++, where the purpose of an if statement is to *do* one thing or another, not to return a value. Also observe that an if expression in our language must have both a then part and an else part, even though only one of the expressions ends up being evaluated.

## Language V3

Language V3 is the same as Language V2 with the addition of a `let` expression. Here are the relevant grammar rules and abstract syntax representations:

```
<exp>:LetExp ::= LET <letDecls> IN <exp>
```

```
LetExp(LetDecls letDecls, Exp exp)
```

```
<letDecls> ::= <VAR> EQUALS <exp>
```

```
LetDecls(List<Token> varList, List<Exp> expList)
```

Notice that we need to change our lexical specification to allow for `LET`, `IN`, and `EQUALS`. Here are the relevant lexical specifications:

```
LET      'let'
IN       'in'
EQUALS   '='
```

Here is an example program in Language V3 that evaluates to 7:

```
let
  three = 2
  four  = 5
in
  +(three, four)
```

The purpose of a `let` expression is to create an environment with variable bindings and to evaluate an expression using these variable bindings.

## Language V3 (continued)

```
<exp>:LetExp ::= LET <letDecls> IN <exp>
               LetExp(LetDecls letDecls, Exp exp)
<letDecls>   **= <VAR> EQUALS <exp>
               LetDecls(List<Token> varList, List<Exp>
```

To evaluate a `LetExp`, we perform the following steps:

0. create a set of local bindings (a `Bindings` object) by binding `<VAR>` symbols to the values of their corresponding `<exp>` expressions in the `<letDecls>` part, where the `<exp>` expressions to the right of `EQUALS` are all evaluated in the enclosing environment;
1. extend the enclosing environment with these local bindings to create a new environment; and
2. use this new environment to evaluate the `<exp>` expression in the `IN` part and return this value as the value of the `letExp` expression.

The `<exp>` part of a `let` expression is called the *body* of the `let` expression.

Examples of `let` expressions are on the next slides, where `=>` means "evaluates to". Remember: *a let expression is an expression, and as such, something!*

## Language V3 (continued)

In an instance of the `LetDecls` class, we require that no `varList` symbol occurs twice. Slide 3.48 shows how we can achieve this.

Now that we can define our own environments, we will remove our environment with bindings for Roman numerals.

In the first example, a new environment is created binding `x` to 3 and `y` to 8. The body of this `let` expression is itself a `let` expression with `z` to 3 and `y` to 8. The environment of the inner `let` extends the environment of the outer `let`, so that the `x` in the expression `+(x,y)` is bound to 3. The expression therefore evaluates to 11.

```
let x = 3 y = 8
in +(x,y) % => 11
```

In the second `let` expression example, a new environment is created binding `x` to 10. The body of this `let` expression is itself a `let` expression with `z` to 3 and `y` to 8. The environment of the inner `let` extends the environment of the outer `let`, so that the `x` in the expression `+(x,y)` is bound to 10. The expression therefore evaluates to 18.

```
let x = 10
in
  let z = 3 y = 8
  in +(x,y) % => 18
```

### Language V3 (continued)

In the third example, two new environments are created. The outer `let` binds `x` to the value of `add1 (x)` and `y` to the value of `add1 (x)`. Both of these `add1 (x)` RHS expressions in the inner `let`, are evaluated in the *outer [enclosing] environment* which has `x` bound to 3. Thus `add1 (x)` evaluates to 4 in both cases. Thus, in the inner environment, `x` is bound to 4 to 4, so that the `+(x, y)` expression evaluates to 8.

```
let x = 3
in
  let
    x = add1 (x)
    y = add1 (x)
  in
    +(x, y)
% => 8
```

Observe also that the `add1` primitive is *not* side-effecting. This expression `add1 (x)` does *not* modify the value bound to `x`: `add1` in this language behaves like `x+1` and *not* like `++x` as you would find in C++ and Java.

### Language V3 (continued)

```
<exp>:LetExp ::= LET <letDecls> IN <exp>
               LetExp(LetDecls letDecls, Exp exp)
<letDecls>   ::= <VAR> EQUALS <exp>
               LetDecls(List<Token> varList, List<Exp> expList)
```

The code for `eval` in the `LetExp` class is straight-forward:

```
LetExp
%%%
    public Val eval(Env env) {
        Env nenv = letDecls.addBindings(env);
        return exp.eval(nenv);
    }
%%%
```

As we show on Slide 3.49, the `addBindings` method returns an `Env` object that extends the `env` environment parameter by adding the bindings given by the `letDecls` declarations. We use this extended environment to evaluate the body expression.



### Language V3 (continued)

```
<exp>:LetExp ::= LET <letDecls> IN <exp>  
                LetExp(LetDecls letDecls, Exp exp)  
<letDecls>    **= <VAR> EQUALS <exp>  
                LetDecls(List<Token> varList, List<Exp>
```

A `LetDecls` object has two instance variables: `varList` is a list of tokens representing the `<VAR>` part of the BNF grammar rule, and `expList` is a list of expressions representing the `<exp>` part of the BNF grammar rule. The reason that these are `Lists` is because the `letDecls` BNF grammar rule says so. Our plan for defining the `addBindings` method in the `LetDecls` class is to evaluate each of the expressions in `expList` in the enclosing environment, binding these values to their corresponding token strings in `varList`, and then use these bindings to extend the enclosing environment given by `env`. Finally, we return this new environment to the `eval` method in the `Rands` class.

The `LetDecls` constructor throws an exception if it finds duplicate tokens in its `varList`. This means that a `let` expression cannot have two identifiers with the same LHS identifier. The code to check for duplicates is inserted into the `LetDecls` class constructor using the (*context-sensitive*) `:init` hook, so the check for duplicates occurs during parsing, not during expression evaluation.

### Language V3 (continued)

By coincidence, the `Rands` object already has an `evalRands` method that evaluates each of the expressions in its `expList` instance variable, so we can use the `Rands` class and its `evalRands` method here.

```
LetDecls  
%%%  
    public Env addBindings(Env env) {  
        Rands rands = new Rands(expList);  
        List<Val> valList = rands.evalRands(env);  
        Bindings bindings = new Bindings(varList, valList);  
        return env.extendEnv(bindings);  
    }  
%%%
```

### Language V3 (continued)

The languages we have discussed do not allow mutation of variables. It might be tempting to think that this Language V3 program is doing so, but it is not. It is not doing so because it is not mutating. It is only evaluating to mutation:

```
let
  x = 3
in
  let
    x = add1(x)
  in
    +(x, x)
```

This program evaluates to 8 (which is not surprising), but in the second `let`, the variable `x` is still bound to 3. To see this, consider the following program:

```
let
  x = 3
in
  +(let x = add1(x) in x, x)
```

The last occurrence of `x` in this expression evaluates to 3 because the inner `let` has scope only through the inner `let` expression. Because the inner `let` expression body, the binding of `x` to 3 remains unchanged. The entire expression evaluates to 7.

### Language V3 (continued)

```
<exp>:LetExp ::= LET <letDecls> IN <exp>
               LetExp(LetDecls letDecls, Exp exp)
<letDecls> ::= <VAR> EQUALS <exp>
               LetDecls(List<Token> varList, List<Exp> expList)
```

Here is another observation you should pay attention to. In the grammar rule, each `<VAR>` symbol is called the *left-hand side* (LHS) of the rule. The corresponding `<exp>` is called its *right-hand side* (RHS). (Don't confuse the LHS and RHS of the grammar rule itself.) All of the RHS expressions `<exp>` are evaluated in the enclosing environment. *The LHS variables become bound to their corresponding RHS expression values after the RHS expressions have been evaluated.* Thus the following expression

```
let p = 4
in
  let
    p = 42
    x = p
  in
    x
```

evaluates to 4.

## Language V4

So far our languages do not allow for anything like repetition. In a based language (ours fall into this category), repetition is typically by recursion, and recursion depends on the ability to apply procedures. So we need to build the capability to define procedures.

In Language V4, we add procedure definitions and procedure application. *procedure* is synonymous with *function*.

Think of a procedure as a “black box” that, when given zero or more inputs, returns a single result value. The number of inputs that a procedure accepts is its *arity*.

To *define* a procedure means to describe how it behaves. To *apply* means to give the procedure the proper number of inputs and to receive the result.

Using mathematical notation, we can *define* a function  $f$  by

$$f(x) = x + 3$$

and we can *apply* the function  $f$  by

$$f(5)$$

The result of this particular application is 8.

## Language V4 (continued)

In Language V4, procedures are treated as values just like integers. To create a `ProcVal` class that extends the `Val` class. This means that a `ProcVal` object can occur anywhere a `Val` object is expected.

Here is an example of a Language V4 program that includes a procedure definition and application.

```
let
  f = proc(x) + (x, 3)
in
  .f(5)
```

In Language V4, a procedure definition starts with the `PROC` token and a procedure application starts with a `DOT`. It is possible that you can define a procedure in one expression, such as

```
.proc(x) + (x, 3) (5)
```

Both of these expressions return the same value, namely the integer 8. The first expression defines a procedure and then applies it, while the second expression defines a procedure and immediately applies it. The first expression returns the integer 8, while the second expression returns a value, but the value is a procedure, not an integer. (One can also define a procedure and then apply it, although this is not a new feature.)

```
proc(x) + (x, 3)
```

also returns a value, but the value is a procedure, not an integer. (One can also define a procedure and then apply it, although this is not a new feature.)

## Language V4 (continued)

Here are some examples of Language V4 programs using procedure

```
let
  f = proc(x,y) +(x,y)
in
  .f(3,8)
  % => 11
```

```
let
  f = proc(z,y) +(10,y)
in
  .f(3,8)
  % => 18
```

```
let x = 10
in
  let
    x = 7
    f = proc(y) +(x,y)
  in
    .f(8)
    % => 18
```

In the third example, the `x` in the `proc` definition refers to the enclosing `let` (where `x` is bound to 10), not to the inner `x` (which is bound to 7). Remember to evaluate the `letDecls`!

## Language V4 (continued)

Now consider the following examples, all of which evaluate to 5:

```
let
  app = proc(f,x) .f(x)
  add2 = proc(y) add1(add1(y))
in
  .app(add2,3)
```

```
let
  app = proc(f,x) .f(x)
in
  .app(proc(y) add1(add1(y)), 3)

.proc(f,x) .f(x) (proc(y) add1(add1(y)), 3)
```

In the first example, observe that we can pass a procedure (in this case `add2`) as a parameter to another procedure. This `app` procedure takes two parameters and returns the result of applying the first actual parameter to the second. The first parameter had better be bound to a procedure for this to work. (If you attempt to apply it throws an exception.)

In the second example, we have eliminated the identifier `add2` and replaced `add2` in the application `.app(add2,3)` with the name of the procedure `proc(y) add1(add1(y))` that used to be called `add2`.

In the third example, we have even eliminated the identifier `app`.

## Language V4 (continued)

Finally consider the following example, which evaluates to 120:

```
let
  fact = proc(f,x)
    if x
    then *(x, .f(f, sub1(x)))
    else 1
in
  .fact(fact, 5)
```

This example, quite a bit more subtle than the previous ones, show achieve recursion – factorial, in this case – using our simple language, not yet support direct recursion!).

One final observation: `add1` is a *primitive*, not a *procedure*. Primitive expressions, and they do not evaluate to anything, so you can't pass an actual parameter to a procedure. In particular, the following will

```
let
  app = proc(f,x) .f(x)
in
  .app(add1, 3)
```

Be aware that the syntax for *applying* a primitive looks somewhat like applying a procedure, except that applying a primitive does *not* use a DOT.

## Language V4 (continued)

We are now prepared to add syntax and semantics to support procedure definition and application and we add grammar rules for procedure definition and application and corresponding abstract syntax classes:

```
<exp>:ProcExp ::= <proc>
                  ProcExp(Proc proc)
<proc>          ::= PROC LPAREN <formals> RPAREN <exp>
                  Proc(Formals formals, Exp exp)
<formals>       ::= <VAR> +COMMA
                  Formals(List<Token> varList)
<exp>:AppExp    ::= DOT <exp> LPAREN <rands> RPAREN
                  AppExp(Exp exp, Rands rands)
```

Before we can go any further, we need to tackle the definition of procedure application, which is what we should get when we evaluate a `ProcExp` expression.

### Language V4 (continued)

A `ProcVal` object must capture the formal parameters as an instance of the `Formals` class, and it must remember its procedure *body* as an instance of the `Exp` class. But what environment should we use to evaluate the procedure body when the procedure is applied? In order to conform to our notion of *static scope*, we must evaluate the procedure body *using the environment in which the procedure was defined*. So any variables in the procedure body which are *not* formal parameters – in other words, the variables that *occur free* in the procedure body – are bound to their values in the environment in which the procedure was defined.

In Programming Languages terminology, the term *closure* refers to an object that captures all of the ingredients necessary to apply a procedure. In our implementation, `ProcVal` objects are closures.

### Language V4 (continued)

The fields of the `ProcVal` class appear here:

```
public class ProcVal extends Val {  
  
    Formals formals;  
    Exp body;  
    Env env;  
  
    public ProcVal(Formals formals, Exp body, Env env) {  
        this.formals = formals;  
        this.body = body;  
        this.env = env;  
    }  
  
    ...  
}
```

Recall that we can do two things with a procedure: *define* it and *apply* it. We will discuss procedure definition shortly, but first we give the semantics of procedure application.

## Language V4 (continued)

```
<exp>:AppExp ::= DOT <exp> LPAREN <rands> RPAREN  
                AppExp(Exp exp, Rands rands)
```

Here are the steps to evaluate a procedure *application* – in other words, an AppExp expression:

0. Evaluate `exp` in the current environment; this must evaluate to an object (a closure) with fields `formals`, `body`, and `env`.
1. Evaluate `rands` (the *actual parameter* [a.k.a. *operand*] expression) in the current environment to get a list of `Vals` (the *arguments*). [Note: this is the same thing when evaluating the `rands` of a `PrimappExp`.
2.
  - a. Create bindings of the procedure's list of formal parameters to the list of values obtained in step 1, and
  - b. use these bindings to extend the environment (`env`) captured by the procedure.
3. Evaluate the `body` of the procedure in the (extended) environment from step 2.

Steps 2 and 3 are carried out by the `apply` method in the `ProcVal` object. The value obtained in step 3 is the value of the AppExp expression evaluated.

## Language V4 (continued)

Let's now examine the detailed semantics of a `ProcExp`, which is a procedure.

```
<exp>:ProcExp ::= <proc>  
                ProcExp(Proc proc)  
  
<proc> ::= PROC LPAREN <formals> RPAREN <exp>  
          Proc(Formals formals, Exp exp)  
  
<formals> ::= <VAR> +COMMA  
            Formals(List<Token> varList)
```

As noted in Slide 3.59, a `ProcVal` closure is constructed with `ins` consisting of the list of formal parameters (a `Formals` object), the procedure body (an `Exp` object), and the environment in which the procedure is defined (an `Env` object).

## Language V4 (continued)

```
<exp>:ProcExp ::= <proc>  
                ProcExp(Proc proc)  
<proc>  
    ::= PROC LPAREN <formals> RPAREN <exp>  
        Proc(Formals formals, Exp exp)  
<formals>  
    **= <VAR> +COMMA  
        Formals(List<Token> varList)
```

The `makeClosure` method in the `Proc` class creates a `ProcVal` in an environment.

```
Proc  
%%%  
    public Val makeClosure(Env env) {  
        return new ProcVal(formals, exp, env);  
    }  
%%%
```

The semantics of the `eval` method in the `ProcExp` class is now tr

```
ProcExp  
%%%  
    public Val eval(Env env) {  
        return proc.makeClosure(env);  
    }  
%%%
```

## Language V4 (continued)

```
<exp>:AppExp ::= DOT <exp> LPAREN <rands> RPAREN  
              AppExp(Exp exp, Rands rands)
```

We provided the structure of the fields in the `ProcVal` class on Slide 3.59. We described the semantics of a procedure application on Slide 3.60. We now give Java code to *implement* application semantics.

We start with the `eval` method in the `AppExp` class. As shown on Slide 3.60, this method carries out steps 0 and 1 of procedure application semantics. On Slide 3.60: it evaluates the `exp` expression – which should evaluate to a `ProcVal` – and then it evaluates the operand expressions to get a list of `Vals`.

It then passes these arguments along to the `apply` method in the `ProcVal` class to carry out steps 2 and 3 of application semantics. This method returns the value of the `AppExp` expression. (As we noted earlier, the operand expressions are the *operands* or *actual parameters*, and their corresponding values are the *arguments*.)

You can find the code on the following two slides.



## Language V4 (continued)

```
<exp>:AppExp ::= DOT <exp> LPAREN <rands> RPAREN
                AppExp(Exp exp, Rands rands)
```

```
AppExp
%%%
```

```
public Val eval(Env env) {
    // evaluate exp in the current environment
    Val v = exp.eval(env); // should be a ProcVal
    // evaluate rands in the current environment
    // to get the arguments (step 1)
    List<Val> args = rands.evalRands(env);
    // let v (step 0) determine what to do next
    v.apply(args, env);
}
%%%
```

Notice that *the operand expressions (the rands) are evaluated in the environment e* *in which the expression is applied*. Also, the current environment *e* is passed as the second parameter to the `apply` method in the `Val` class, even though we see that the `apply` method in the `ProcVal` class does not actually

## Language V4 (continued)

The only thing we have left is to implement the behavior of the `apply` method in the `Val` class. Since we want `apply` only to be meaningful for a `ProcVal`, we define a default behavior in the (abstract) `Val` class to throw an exception for anything but a `ProcVal`:

```
public Val apply(List<Val> args, Env e) {
    throw new PLCCEException("Cannot apply " + this);
}
```

For a `ProcVal`, here's the implementation of `apply`. Notice that the `ProcVal` implementation carries out steps 2a, 2b, and 3 in the semantics for evaluating application (Slide 3.60).

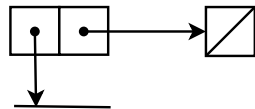
```
public Val apply(List<Val> args, Env e) {
    // bind the formals to the arguments (step 2a)
    Bindings bindings = new Bindings(formals.varList(), args);
    // extend the captured environment with these bindings
    Env nenv = env.extendEnv(bindings);
    // and evaluate the body in this new environment
    return body.eval(nenv);
}
```

Language V4 also checks for duplicate identifiers while parsing the input. This code is inserted into the `Formals` class using the `:init` hook.

## Drawing Environments

On Slide 2.20, we showed how to display environments as a linked list. In the list is a pair (`EnvNode`) consisting of a reference to a `Bir` (which we have called *local bindings*) and a reference to the next environment. The end of the list is an empty environment, an `EnvNull` object, which is a box with a slash through it. We usually display the linked list *horizontally*, with the head of the list at the left and the empty environment at the right. To display the local bindings as an array (it's actually an `ArrayList`) stacked vertically. Each binding is a pair consisting of an identifier and a value.

The *initial environment* in Language V4 is a linked list consisting of one `EnvNode` with an empty local environment (no bindings) and a reference to an `EnvNull` object. Here is how we display the initial environment:



**To simplify things in Languages V4 and V5, we omit displaying the empty local environment, so we display the initial environment as follows:**



## Drawing Environments (continued)

There are exactly two ways in which programs in Language V4 create new environments using the `extendEnv` method:

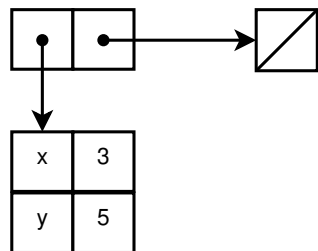
- evaluating a `let` expression
- evaluating a procedure application

A `let` expression creates a list of local bindings: each binding uses the identifier as its `id` field and the value of the RHS expression as its `val` field. The RHS expressions are evaluated *in the enclosing environment* that the RHS expressions are evaluated *in the enclosing environment* being created. An example expression is given on the next slide.

### Drawing Environments (continued)

```
let
  x=3
  y=5
in
  +(x,y)
```

This `let` expression creates an environment that extends the initial environment with bindings for `x` and `y`. This extended environment is then the body expression `+(x,y)` is evaluated. The environment diagram extended environment is shown here:

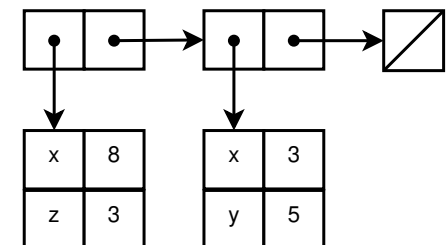


### Drawing Environments (continued)

Now consider the following expression, with nested `lets`. The inner `let` extends the environment defined by the outer `let` (with one node, as shown on the previous page), so the environment of the inner `let` is a linked list with two nodes:

```
let
  x = 3
  y = 5
in
  let
    x = +(x,y) % the RHS evaluates to 8
    z = x      % the RHS evaluates to 3 (why?)
  in
    +(x,y)
```

In the following diagram, the leftmost node is the environment created by the outer `let`:



The inner `let` body expression evaluates to 13 (why?).

## Drawing Environments (continued)

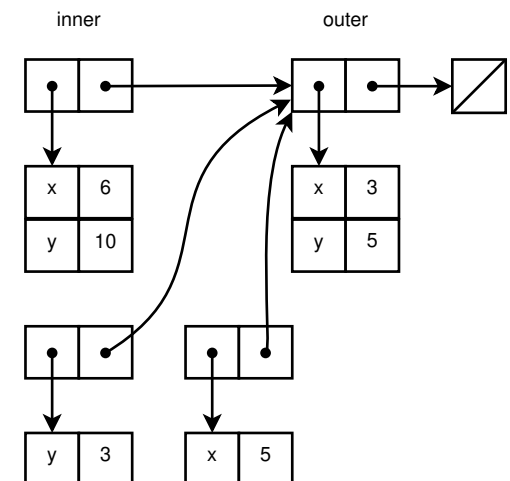
Since a `let` expression is an expression, it must evaluate to a value. An expression can occur as the RHS of a binding in another `let` expression. This example:

```
let % outer
  x = 3
  y = 5
in
  let % inner
    x = let y=x in +(x,y) % LHS x is bound to 6
    y = let x=y in +(x,y) % LHS y is bound to 10
  in
    +(x,y) % evaluates to 16
```

The environment defined by the outer `let` has one node with bindings (to 3 and 5, respectively). The inner `let` extends the environment of the outer `let`. The inner environment has two nodes. The RHS expressions for the inner `let` are evaluated in the environment defined by the outer `let`. These RHS expressions are themselves `let` expressions, each of which extends the environment defined by the outer `let`. A total of four environments are created: the outer `let` (extending the initial null environment), the inner `let` (extending the outer `let`), and one for each of the RHS expressions in the inner `let` (extending the inner `let`). The next slide shows all of these environments.

## Drawing Environments (continued)

```
let % outer
  x = 3
  y = 5
in
  let % inner
    x = let y=x in +(x,y) % LHS x is bound to 6
    y = let x=y in +(x,y) % LHS y is bound to 10
  in
    +(x,y) % evaluates to 16
```



## Drawing Environments (continued)

In the definition of the `ProcVal` class, a `ProcVal` object has three

```
public Formals formals; // list of formal parameters
public Exp body;        // procedure body
public Env env;         // captured environment
```

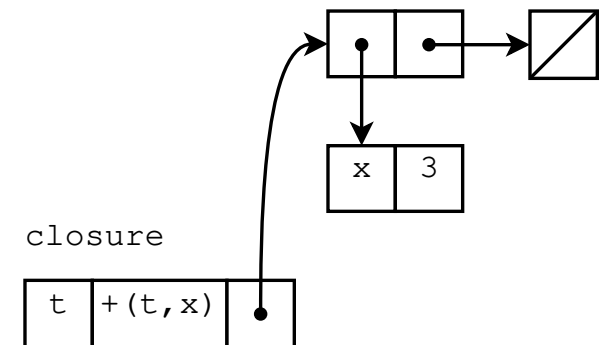
Here, the *captured environment* is the environment in which the procedure was defined. For example, consider the following expression:

```
let
  x = 3
in
  proc(t) +(t,x)
```

This expression evaluates to a `ProcVal`: its `formals` field consists of a single string, 't', its `body` is the expression '+ (t, x)', and its `env` is the one defined by the `let`, having a single binding of `x` to 3. (Recall that we also use the term *closure* to refer to a `ProcVal` object.) We normally display a `ProcVal` object as a rectangle with its three components in this order: formals, body, and captured environment. We show the captured environment (`env`) as an arrow pointing to the appropriate environment in which the procedure definition occurs. The following diagram shows the `ProcVal` that results from the evaluation of the above expression:

## Drawing Environments (continued)

```
let
  x = 3
in
  proc(t) +(t,x)
```



## Drawing Environments (continued)

We described the rules for *applying* a procedure on Slide 3.60. In the `ProcVal` class, the key steps are: (2a) bind the formal parameters to the values of the actual parameter expressions (object); (2b) extend the captured environment with these bindings to the environment; (3) evaluate the procedure body using the new environment. The value obtained in step (3) is the value of the procedure application.

Consider the following expression, which is the same as the previous one except that we apply the procedure to the actual parameter 5:

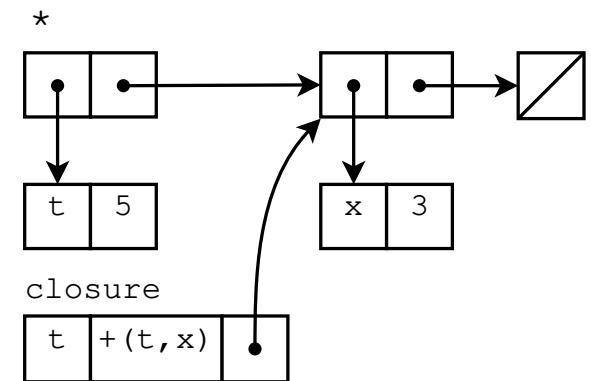
```
let
  x = 3
in
  .proc(t) +(t,x) (5)
```

From the above discussion, this procedure application creates a local formal parameter `t` to the value 5 (the actual parameter expression binding is used to extend the environment captured by the procedure). The extended environment is used to evaluate the body of the procedure. The value of the procedure application is 8.

The following page displays the environment created by this application with an asterisk '\*'.  
with an asterisk '\*'.

## Drawing Environments (continued)

```
let
  x = 3
in
  .proc(t) +(t,x) (5)
```



## Drawing Environments (continued)


Finally, consider this example:

```
let
  x = 3
in
  let
    f = proc(t) +(t,x)
  in
    .f(5)
```

The value of this expression is also 8. Evaluating this expression creates environments: one with a binding of `x` to 3, another with a binding of `f` and a third created by applying `f` to the argument 5. The resulting diagram is shown on the following page. The environment in which the body of `f` is marked with an asterisk `*`.

## Drawing Environments (continued)

```
let
  x = 3
in
  let
    f = proc(t) +(t,x)
  in
    .f(5)
```



## Language V4 (continued)

Once we have procedures, we can entirely eliminate the `let` construct. For example:

```
let
  p = 3
  q = 5
in
  +(p, q)
```

This can be re-written as an application of an anonymous (un-named) procedure as follows:

```
.proc(p, q) +(p, q) (3, 5)
```

## Language V4 (continued)

In general, a `let` expression

```
let
  v1 = e1
  v2 = e2
  ...
in
  e
```

can be re-written as an equivalent procedure application expression

```
.proc(v1, v2, ...) e (e1, e2, ...)
```

This conversion is *algorithmic*: it can be carried out unambiguously. Should we ditch the `let` construct? The reason is simple: it's easier to think about a program with a `let` in it than one without. The `let` construct aligns the LHS variables `v1`, `v2`, *etc.* physically close to their corresponding RHS expressions `e1`, `e2`, *etc.* it's cognitively easy for the reader to see how these LHS variables are bound to the values of their RHS expressions. In the equivalent procedure application, the formal parameters `v1`, `v2`, *etc.* are physically distant from their corresponding RHS expressions, making it difficult to visualize these bindings. A `let` expression is an example of *syntactic sugar*: a syntactic and semantic construct that is not needed to express an equivalent way of expressing it in the language but that programmers find easier to read, understand, and use.



## Language V4 (continued)

Though Language V4 does not support direct recursion, its support for first-class entities – that is, they are values that are treated the same as values, so they can be passed as parameters and returned as values – is as good as direct recursion. Here is another example that recursively computes the factorial using an “accumulator” and tail recursion (we will return to this topic later).

```
let
  fact = proc(x)
    let
      factx = proc(f, x, acc)
        if zero?(x)
          then acc
          else .f(f, sub1(x), *(x, acc))
      in
        .factx(factx, x, 1)
    in
      .fact(5)
```

Observe that the identifier `f` that appears in the `proc(f, x, acc)` is a *formal parameter* name that binds all occurrences of `f` that appear in the body. You may find it instructive to display all of the environments during the evaluation of this expression. (Replace ‘5’ by ‘2’ to make

## Language V4 (continued)

Finally, Language V4 includes the ability to evaluate a sequence of expressions, returning the value of the last expression. The component expressions of a sequence expression are always evaluated left-to-right. Sequence expressions are particularly useful now because our language is not *side-effecting*, but they will turn out to be useful in later languages that do support side-effects.

```
<exp>:SeqExp ::= LBRACE <exp> <seqExps> RBRACE
                SeqExp(Exp exp, SeqExps seqExps)
<seqExps> ::= SEMI <exp>
                SeqExps(List<Exp> expList)
```

The semantics of evaluating a `SeqExp` are given here:

```
SeqExp
%%%
    public Val eval(env) {
        Val v = exp.eval(env);
        for (Exp e : seqExps.expList)
            v = e.eval(env);
        return v;
    }
%%%
```

## Language V4 (continued)

Observe that we evaluate every expression in the list but only return the last one:

```
{1; 3; 5}  
% => 5
```

```
{42}  
% => 42
```

We can use the sequence construct to enclose a single expression though this looks a bit unwieldy. Here's an example:

```
.{proc(t,u) +(t,u)} (3,4)
```

The braces in this expression are not required, but they may help to the extent of the `proc` definition. Don't get into the habit of doing this, however: throwing in extra braces can result in an expression that is unnecessarily *noisy* and that is actually more difficult to read than one without them.

## Language V5

We normally prefer to use direct recursion instead of using the (controllable) tricks on Slides 3.56 and 3.80. For example, we would like to

```
let  
  fact = proc(x) if zero?(x) then 1 else *(x, .fact(x-1))  
in  
  .fact(5)
```

But this does not work! Why??

Remember that in a `let`, the RHS expressions (the expressions to the right of the '=' tokens) are all evaluated in the environment that encloses the `let`. Only after all the RHS expressions have been evaluated do we bind each of the identifiers to their RHS values.

In the definition of the `proc` above, the `proc` body refers to the identifier `fact`, which is free in the procedure definition – there is no `fact` in its environment. Thus an attempt to apply the `proc` fails because of an unbound identifier.

## Language V5 (continued)

In order to solve this problem, we create a new `let`-like environment for direct recursion. Called `letrec`, it allows us to define procedures with direct recursion.

This is what we want:

```
letrec
  fact = proc(x) if zero?(x) then 1 else *(x, .fact(x-1))
in
  .fact(5)
% => 120
```

## Language V5 (continued)

Here is the grammar rule and associated abstract syntax class:

```
<exp>:LetrecExp ::= LETREC <letDecls> IN <exp>
LetrecExp(LetDecls letDecls, Exp exp)
```

The RHS expressions in a `letrec` are evaluated in the order in which they appear, using an environment where *all* of the previous (LHS, RHS) bindings in the `letrec` are accessible. In addition, if the RHS is a procedure, the entire environment created by *all* of the (LHS, RHS) bindings in the `letrec` means that procedures defined in a `letrec` can refer to each other, and they can call themselves recursively.

This is unlike a normal `let`, in which the RHS expressions are all evaluated in the *enclosing* environment, and the (LHS, RHS) bindings created by the RHS are not accessible in the body of the `let`. Moreover, in a `let`, the RHS expressions can be evaluated in any order, which is sometimes called *parallel evaluation*.

Notice that the syntax of a `letrec` is the same as the syntax of a `let`. The difference between the semantics of `let` and `letrec` is in the way they build the environment in which the `letDecls` bindings are created.

We proceed to describe how `letrec` evaluation is handled.

## Language V5 (continued)

To implement the recursive behavior of a `letrec` as described above, we will add a new method `addLetrecBindings` in the `LetDecls` class. This method is passed the environment in which the `letrec` expression appears and returns a new environment as described in the following steps.

0. Extend the environment `actual` parameter with an empty `Bindings` object of sufficient size to hold all of the variable bindings. Assign this new environment to the `env` parameter.
1. The two fields in the `LetDecls` class are `List<Token> varList` and `List<Exp> expList`. Create iterators for these two lists and iterate over them together, in order. For each step in the iteration, get the next token from the `varList`, and save its `String` representation in a variable `str`. Also, get the next expression `exp` from the `expList`, evaluate it in the environment `env` obtained in Step 0, and save its value in a variable `val`. Create a new `Binding(str, val)` and add it to the `env` environment obtained in Step 0. This binding now becomes part of the local bindings of the environment together with the other local bindings previously added during the iteration.
2. Once all of the new bindings have been added to `env`, return `env` as the result of this method.

## Language V5 (continued)

The implementation of `addLetrecBindings` in the `LetDecls` class is shown here:

```
LetDecls
%%%
    public Env addLetrecBindings(Env env) {
        // Step 0
        env = env.extendEnv(new Bindings(varList.size(), null));
        // Step 1
        Iterator<Token> varIter = varList.iterator();
        Iterator<Exp> expIter = expList.iterator();
        while (varIter.hasNext()) {
            String str = varIter.next().toString();
            Val val = expIter.next().eval(env);
            env.add(new Binding(str, val));
        }
        return env; // Step 2
    }
%%%
```

Notice that we have previously defined an `addBindings` method in the `LetDecls` class (see Language V3) used to implement the evaluation of a `let` expression. The `addLetrecBindings` method simply builds on top of the `addBindings` method.

## Language V5 (continued)

```
<exp>:LetrecExp ::= LETREC <letDecls> IN <exp>  
LetrecExp(LetDecls letDecls, Exp ex
```

Recall that the `LetDecls` constructor checks for duplicate LHS id parsing. Since the `LetrecExp` grammar rule uses `LetDecls`, expression also makes this check.

We can now evaluate a `LetrecExp` object in exactly the same way object:

```
LetrecExp  
%%  
public Val eval(Env env) {  
    Env nenv = letDecls.addLetrecBindings(env);  
    return exp.eval(nenv);  
}  
%%
```

The principal idea, then, is to evaluate the RHS expressions of a environment that (self-referentially) includes all of the bindings in t

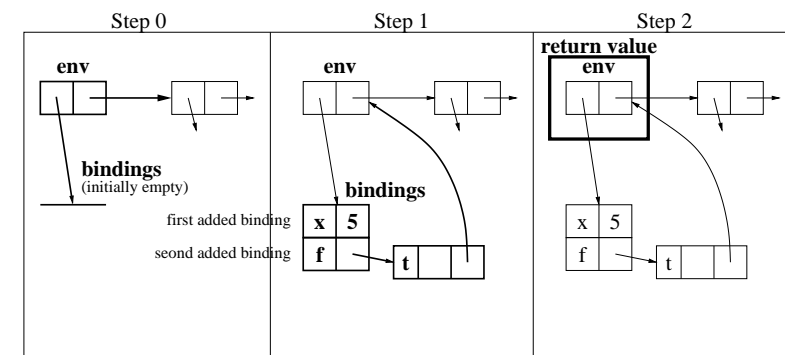
## Language V5 (continued)

This picture illustrates the three steps carried out in `addLetrecBindings`, for the following 1

```
letrec  
  x = 5  
  f = proc(t) *(t,x)  
in  
  .f(42)
```

0. Create an extended environment by extending the enclosing environment with an empty list
1. In the order in which the LHS identifiers appear, create a `Binding` of the LHS identifier (for example) to the value of its corresponding RHS expression (5 and then `proc(t) ...` in the example) – and add this binding to the extended environment
2. Once all of the (LHS, RHS) bindings have been added to the extended environment, return the extended environment as the value of `addLetrecBindings`.

Observe that the environment captured by procedure `f` knows about the binding of `x` to 5, so the procedure evaluates to 210. If this had been a `let` instead of a `letrec`, the `x` in the body of `f` would be un



## Language V5 (continued)

The `letrec` construct allows us to define *mutually recursive procedures* or more procedures that call each other in a recursive fashion. For example:

```
letrec
  even? = proc(x) if zero?(x) then 1 else .odd?(sub1 x)
  odd? = proc(x) if zero?(x) then 0 else .even?(sub1 x)
in
  .even?(11) % => 0 (false)
```

**[Exercise (not to hand in):** See if you can define the `odd?()`/`even?` recursive procedures in Language V4 without `letrec`.]

Notice that we have used `?` in the variable names for the `even?` procedures to suggest that these procedures should be considered as returning true (1) or false (0). This is a lexical feature we have added to and beyond.

## Language V6

So far, our source language has no capability to define top-level variables that persist from one expression evaluation to another. A *top-level* variable is a binding in the initial (“top-level”) environment. We would like to extend our languages to allow for such definitions. All we need to do is to add the *initial environment*, the environment that all expressions in the source language start out with.

The initial environment for our languages is the static `Env` object created in the `Program` class, obtained by calling the `initEnv()` static method of the `Env` class. Notice that this initial (top-level) environment starts out having no bindings.

Our strategy for making top-level definitions is to take advantage of the `addBinding()` method in the `Env` class. To add a new top-level definition, we create a `Binding` object and add it to the top-level `Env` object. Once we add bindings in this way, the bindings will be known in any subsequent expression evaluation that starts in the initial environment.

## Language V6 (continued)

Since a “program” can now have two forms – a top-level “define” or “eval” for evaluation, we need to have two grammar rules for the <program> non-terminal. Here are their grammar rules and corresponding abstract syntax classes:

```
<program>:Define ::= DEFINE <VAR> EQUALS <exp>  
Define(Token var, Exp exp)  
<program>:Eval ::= <exp>  
Eval(Exp exp)
```

## Language V6 (continued)

Here is an example of expressions that use the define feature in the language:

```
define i = 1  
define ii = add1(i)  
define iii = add1(ii)  
define v = 5  
define x = 10  
define f = proc(x) if zero?(x) then 1 else *(x, .f(.f(v)  
    .f(v) % ERROR: g is unbound  
define g = proc(x) sub1(x)  
    .f(v) % => 120 -- g is now bound  
    .f(iii) % => 6
```

As long as you stay in the Rep loop, the defined variable bindings are remembered.

Notice that, in the definition for `f`, the body of the procedure refers to a variable named `g`, but `g` hasn't been defined yet. The attempt, in the next line, to evaluate `.f(v)` fails. After defining `g` on the following line, evaluating `.f(v)` succeeds. This is because by the time you attempt to apply `f` the second time, the `g` variable has been defined, and the body of `f` now recognizes its definition.

## Language V6 (continued)

Notice that for top-level procedure definitions, `define` works similarly in terms of being able to support direct recursion. This is because procedure definition captures (in a closure) the initial environment, verified every time another top-level definition is encountered. When binding to the top-level environment, the binding gets added to the instead of extending the top-level environment. In this way, all of the closures can access this binding, as well as any others that may crop up following works:

```
define even? = proc(x)
  if zero?(x) then 1 else .odd?(sub1(x))
  .even?(11) % => Error: unbound procedure odd?
define odd? = proc(x)
  if zero?(x) then 0 else .even?(sub1(x))
  .even?(11) % => 0
  .odd?(11) % => 1
```

Observe that a top-level `define` can *redefine* a previous definition, looking up the LHS identifier in the top-level environment. If an identifier already exists in the top-level environment, we replace the binding with the value of the new RHS.

## Language V6 (continued)

Since a `define` evaluates its RHS in the *current* environment, we can (save) the value of a variable using a `let`, even though a subsequent `define` redefines the variable. Consider this example:

```
define x = 2
define f = proc() x % x is the top-level x
  .f() % evaluates to 2
define x = 3 % redefine top-level x
  .f() % now evaluates to 3
```

Compare this to the following:

```
define x = 2
define f =
  let
    x = x % the RHS is the current environment
    % and the LHS is a local binding

  in
    proc() x % the proc captures the current environment
    .f() % evaluates to 2
define x = 3 % redefine top-level x
  .f() % local copy still evaluates to 2
```