

Copyright (C) 2020 Timothy Fossum

Permission is granted to copy, distribute and/or modify this document under the terms of the GNU Free Documentation License, Version 1.1 or any later version published by the Free Software Foundation. A copy of the license is included in the file "COPYING", entitled "GNU Free Documentation License".

Compared to most object-oriented programming languages such as Java and C++, our classes and objects are *first class* – that is, they can be created at any time and in any environment, they can be passed as parameters to procedures, and they can be returned as values of procedure applications.

A *class* is an expressed value that captures a collection of variables (called *fields*) and procedures (called *methods*) and that serves as a template to create expressed values called *instances* of the class. The instances are also called *objects*, to distinguish them from classes. You can think of a class as a factory that identifies how to churn out an arbitrary number of instances of the class.

All classes of an OBJ program belong to a *class hierarchy*, which is a forest – a set of trees – with an unnamed class at the root of each of its trees, and with program-created classes at the other nodes of the trees.

In the class hierarchy, a class X that occurs as a child node of class Y in the class hierarchy is called a *subclass* of Y , and Y is called a *superclass* of X .

When an object of class X is instantiated, instances of each of the classes that lie on the path from X to the root of the tree are created, and the combination of all those instances is considered as the resulting object.

If Y is the superclass of X , then an object x created from class X “contains” an object y created from Y . The object y is called the *parent* of x , and likewise x is called the *child* of y .

A class may also have `static` variables whose values are shared among all instances of the class.

Consider the following example.

```
define c1 =  
  class % extends an unnamed class  
    field x  
    field y  
  end  
define c2 =  
  class extends c1  
    field z  
    field x  
  end  
define obj1 = new c1  
define obj2 = new c2  
<obj1>set x = 3  
<obj2>set x = 5  
<obj1>x % evaluates to 3
```

In this example, `c1` is a class and `c2` is a subclass of `c1`. Object `obj1` is an instance of `c1` and `obj2` is an instance of `c2`.

We define the concrete and abstract syntax of classes and objects. We also add the additional reserved word `nil` to our language, which represents a separate expressed value not shared by any other data type. When used in conditional expressions, `nil` is considered false.

```
<exp>:NilExp      ::= NIL  
                    NilExp()  
<exp>:ClassExp    ::= CLASS <ext> <statics> <fields> <methods> END  
                    ClassExp(Ext ext, Statics statics, Fields fields, Methods methods)  
<ext>:Ext0        ::=  
                    Ext0()  
<ext>:Ext1        ::= EXTENDS <exp>  
                    Ext1(Exp exp)  
<statics>        **= STATIC <VAR> EQUALS <exp>  
                    Statics(List<Token> varList, List<Exp> expList)  
<fields>         **= FIELD <VAR>  
                    Fields(List<Token> varList)  
<methods>        **= METHOD <VAR> EQUALS <proc>  
                    Methods(List<Token> varList, List<Proc> procList)  
<exp>:NewExp      ::= NEW <exp>  
                    NewExp(Exp exp)  
<exp>:EnvExp      ::= LANGLE <exp>vExp RANGLE <exp>eExp  
                    EnvExp(Exp vExp, Exp eExp)
```

Every `Exp` expression evaluates to a `Val` object, and a `ClassExp` is no exception, so evaluating such a `ClassExp` (*i.e.*, calling its `eval` method) returns a `Val` object. Looking only at the syntax, it seems reasonable that such an object has a superclass (the `ext` [for `EXTENDS`] part), a set of static identifiers and associated values, a set of field names, and a set of method names and associated procedures.

We define a `StdClass` object in the file `class`, which extends the `Val` class. (Actually, it extends `ClassVal` class which in turn extends the `Val` class.) Evaluating a `ClassExp` expression returns an instance of `StdClass`.

[Warning: In this chapter we implement classes and objects in our defined language (OBJ), using classes and objects in our defining language (Java). This can be confusing. For example, a `ClassExp` evaluates to an object in Java that represents a class in our defined language. Be sure that you have a clear understanding of the context in which the terms “class” and “object” are being used in the following discussion.]

A `StdClass` object (we're in the defining language, Java, now) has the following instance variables:

```
public ClassVal superClass;  
public Bindings staticBindings;  
public Fields fields;  
public Methods methods;  
public Env staticEnv;
```

The `superClass` variable is a reference to a `ClassVal` object (defining) that is the superclass (defined) of this class (defined). The `staticBindings` variable is a reference to the the bindings of the static variable names (defined) to their RHS values (defined). The RHS values are evaluated in the current static environment. (More about this later.) [OK, so I'll stop the “(defined)” and “(defining)” stuff now, but *pay attention*.]

The `fields` and `methods` variables are references to the `Fields` and `Methods` objects that capture the class field names and method names and procedures – these are not evaluated yet.

The `staticEnv` variable is a reference to the static environment of this class: it starts out extending the static environment of the superclass with an empty `Bindings` object named `staticBindings`. New bindings are added to the `staticBindings` list using the variable definitions in the `statics` parameter of the `StdClass` constructor. Each static LHS identifier is bound to the value of its RHS, where the RHS expression is evaluated in the current static environment. These bindings are created in order (first to last) as in top-level `defines`.

Two predefined identifiers are inserted initially into the list of static bindings: one binds the identifier `myclass` to (a reference to) this class itself, and another binds the identifier `superclass` to (a reference to) the superclass of this class.

If the class expression does not have an `extends` part, its superclass defaults to an unnamed “parentless” class (an `EnvClass` object in the Java implementation) whose static environment is the environment in which the class is created. In this way, all of the RHS expressions in the `static` definitions of such a class have access to other static bindings, including the bindings in the environment in which the class expression appears.

A class expression that specifies an explicit superclass has an environment that extends the static environment of that superclass, and – going up the superclass chain – has access to the bindings in all of the superclass static environments.

If a class expression has a static definition for a variable that also appears as a static variable in a superclass, that definition shadows the superclass variable. We disallow duplicate LHS variable names in a given class expression’s static definitions, including redefinitions for `myclass` or `superclass`, although our implementation does not check for this.

Although counter-intuitive, objects are actually simpler than classes, because an object is essentially an environment!

An `ObjectVal` is a Java class that extends the `Val` class. It has a single instance variable:

```
public Env objectEnv;
```

The `new` operator in our defined language takes a class expression and returns a Java `ObjectVal` instance that essentially couples the `static` bindings of the class with bindings for the class fields and methods. Since our language does not define an explicit constructor in class expressions, object fields are initialized to `nil`. The method variable names in the class definition are bound to procedure closures that capture the environment that includes bindings for the class `static` variables (from the `staticBindings` field), along with bindings for the fields as described above. The method closures are created as in `letrec`, so they can refer to themselves recursively. As with static variables, we disallow duplicate method variable names, although our implementation does not check for this.

Before an object builds the environment that is specific to the fields and methods of the class, it creates an instance of the superclass of the class and adopts the environment of the superclass object (bound to the variable `super`) before adding `static`, `field`, and `method` bindings. Since creating the superclass object may itself involve creating an instance of *its* superclass, object creation continues up the class hierarchy until a parentless class is found, at which point there is no further `super` object to create.

At the top of the chain of `super` objects, the identifier `self` is bound to a reference to the object being created. In this way, the object can refer to itself through the `self` identifier. (In Java, we call it `this` instead of `self`.) Methods declared in superclasses that refer to `self` will “see” the original object, allowing for dynamic dispatch of method calls, an important feature of object-oriented languages. We call this a *deep* binding.

As objects are created up the superclass chain, the identifier `this` is bound to the object created by that class. We call this a *shallow* binding.

Notice that an object can see all of the `static` bindings up the superclass chain, but that if a `static` variable is bound to a procedure (or some other value that captures an environment), the procedure captures only the static environment of the class and cannot “see” any of the fields or methods – including the `self` and `this` identifiers – in its environment.

A parentless class is an instance of the Java `EnvClass` class, which has one instance variable:

```
public Env staticEnv;
```

Here is the code for the constructor, which is called when processing the `<ext>:Ext0` grammar rule. Notice that the static superclass identifier of a parentless class is bound to `nil`.

```
public EnvClass(Env env) {  
    Bindings staticBindings = new Bindings();  
    // the static environment of this class extends the current env  
    staticEnv = env.extendEnvRef(staticBindings);  
    // create bindings for these static symbols ...  
    staticBindings.add("superclass", new ValRef(Val.nil));  
    staticBindings.add("myclass", new ValRef(this));  
}
```

As noted earlier, the static environment of a parentless class is the environment in which it is created, which is obtained from the `Env env` parameter in the `EnvClass` constructor.

A standard class (an instance of the Java `StdClass` class) has a similar constructor, except that it builds on the environment of its superclass as described earlier. The code for this constructor is in the file `class`.

As described on slides 9 and 10, as an object is created, the `self` identifier is bound to a reference to the object in the `makeObject` method in the parentless class. But how can this binding take place when the object has not yet been completely created? We do this by creating a Java `ValRef` object named `objRef`, with a dummy initial value (`nil`, to be precise), and pass this to the `makeObject` method.

The `objRef` is passed up the superclass chain through successive `makeObject` calls. When `makeObject` reaches a parentless class, the `self` identifier is put into the object environment, bound to `objRef`. After the original object has been completely created – that is, after all of the `makeObject` calls have returned, `objRef` is rebound to the newly created object with a call to `setRef`. Here is the complete `eval` code for a new expression in the `NewExp` part of the code file:

```
public Val eval(Env env) {
    // get the class from which this object is created
    Val val = exp.eval(env);
    // create a reference to a dummy value (nil)
    Ref objRef = new ValRef(Val.nil);
    // let the class create the object
    ObjectVal objVal = val.makeObject(objRef);
    // set the reference to the newly created object
    return objRef.setRef(objVal);
}
```

The `makeObject` method is defined for both a `StdClass` and an `EnvClass`. (For all other `Val` objects, `makeObject` throws an exception.)

In `StdClass`, `makeObject` first creates an instance of the superclass, and then stitches together an environment that includes the static bindings, the fields, and the methods. Three fields are created and initialized automatically: `self` is bound to the base object (a deep binding), `super` is bound to the instance of the superclass, and `this` is bound to this object (a shallow binding). The remaining named fields are bound to `nil`, and the methods are bound to closures as in `letrec`. The code for `makeObject` is given on the next two slides.

```
public ObjectVal makeObject(Ref objRef) {
    // create the parent object (recursively)
    ObjectVal parent = superClass.makeObject(objRef);
    // this object's environment extends the parent object's env
    Env e = parent.objectEnv;
    // add this class's static bindings (including myclass, etc.)
    e = e.extendEnvRef(staticBindings);
    // bind 'super' to the parent object (deep)
    // bind 'self' to the base object
    // bind 'this' to this object (shallow)
    // 'self' is unnecessary here, except that it speeds up lookups
    Bindings fieldBindings = new Bindings();
    e = e.extendEnvRef(fieldBindings);
    fieldBindings.add("super", new ValRef(parent)); // parent object
    fieldBindings.add("self", objRef); // deep
    fieldBindings.add("this", new ValRef(objectVal)); // shallow
    // next bind all of this object's instance fields to nil
    for (Token t : fields.varList) {
        String s = t.toString();
        fieldBindings.add(s, new ValRef(Val.nil));
    }
    e = e.extendEnvRef(fieldBindings);
}
```

... continued on next slide ...

... continued from previous slide ...

```
// bind all this object's methods as in letrec
if (methods.varList.size() > 0) {
    LetrecDecls methodDecls =
        new LetrecDecls(methods.varList, methods.procList);
    e = methodDecls.addBindings(e);
}

// create the object and return it
return new ObjectVal(e);
}
```

The `makeObject` method for `EnvClass` is simple, since it's always the last superclass object that needs to be constructed. It extends the environment in which the `EnvClass` is created (*not* the top-level environment) with a single field binding of `self` to (a reference to) the object being created.

```
public ObjectVal makeObject(Ref objRef) {  
    // start with the env. in which the EnvClass is defined  
    Env e = staticEnv;  
    // add the field binding 'self' to refer to this object  
    Bindings fieldBindings = new Bindings();  
    fieldBindings.add("self", objRef);  
    e = e.extendEnvRef(fieldBindings);  
    return new ObjectVal(e);  
}
```

Such an `ObjectVal` can access the environment in which the parentless `EnvClass` object is created, which is always the environment in which a standard class is created with no `extends` part.

Observe that the RHS expressions in static definitions are evaluated in the static environment of the class, which extends the static environment of its superclass (and so on...). For example, consider the following expression:

```
let
  c = class static x = 5 end
  x = 3
in
  class extends c
    static y = x
  end
```

In this case, `y` is bound to 5, not 3.

The $\langle \dots \rangle$ operator, when applied to a class, extracts the static environment of the class. This can be used to find the binding of a variable in that environment (as in $\langle c \rangle x$) or to apply a procedure (as in $\cdot \langle c \rangle f (\dots)$).

When making a procedure application such as $\cdot \langle c \rangle f (\dots)$, it is important to know

- the environment in which f is evaluated, and
- the environment in which the actual parameters to f are evaluated.

To evaluate the expression $.<c>f(. . .)$, we evaluate the procedure f in the environment determined by $<c>$, whereas we evaluate the the actual parameter expressions in the environment in which the application $.<c>f(. . .)$ is made. For example, in

```
define f = proc(t) *(2,t)
define c =
  class
    static x = 3
    static f = proc(t) t
  end
let
  x = 5
in
  .<c>f(x)
```

the expression $.<c>f(x)$ evaluates to 5, because f is bound to the procedure given in c , but its actual parameter x evaluates to 5, since the actual parameter expression is evaluated in the local `let` environment.

Keep in mind that the following two expressions are not equivalent:

$$.<c>f (. . .)$$
$$<c>.f (. . .)$$

In the first expression, f is evaluated in the static environment of c . Then the procedure bound to f is applied to the actual parameters $(. . .)$ which are evaluated in the current environment, *not* in the static environment of c .

In the second expression, the entire expression $.f (. . .)$ is evaluated in the static environment of c , which means that the actual parameters $(. . .)$ are also evaluated in this static environment.

In the example on the previous slide, if the final expression was $<c>.f (x)$ instead of $.<c>f (x)$, it would evaluate to 3, since x is bound to 3 in the static environment of c .

To clarify, the above two expressions can be re-written as follows:

$$.\{<c>f\} (. . .)$$
$$<c>\{.f (. . .)\}$$

The special operator @ returns an object that captures the *current* environment, whatever that may be. (Recall your homework assignment that introduced this notation.) This operator may be used to pass a captured environment as an argument to a procedure application or to assign it to a variable for later reference.

```
AtExp
```

```
%%%
```

```
public Val eval (Env env) {  
    return new ObjectVal (env) ;  
}
```

```
%%%
```

The special operator @@ does the same thing, except that it also displays the current environment in a human-readable way.

Some examples:

```
define x = 11
define y = 42
define z = 666
define xyenv =
  let
    x = 3
    y = 5
  in
    @
```

```
<@>x      % => 11
```

```
<@>y      % => 42
```

```
<@>z      % => 666
```

```
<xyenv>x  % => 3
```

```
<xyenv>y  % => 5
```

```
<xyenv>z  % => 666 (the 'let' extends the top-level env)
```


Observe that for any expression exp , the following two expressions evaluate to the same values:

$\langle @ \rangle \text{exp}$

exp

Unlike the `new` operator in Java, Our `new` operator does not take any arguments, and all of the fields are initialized to `nil`. We can, of course, initialize fields by calling a method. Here's an example:

```
let
  c = class
    field x
    field y
    method init = proc(a,b) {set x=a; set y=b ; self}
  end
in
  let
    o = .<new c>init(3,4)
  in
    <o>+(x,y) % => 7
```

Since the `init` method returns `self`, the value of the expression `.<new c>init(3,4)` is the same object as the one created by the call to `new`, with its fields `x` and `y` set to values 3 and 4, respectively.

The last slide's code is repeated here (in slightly compressed form):

```
let
  c = class
    field x
    field y
    method init = proc(a,b) {set x=a; set y=b ; self}
  end
in
  let o = .<new c>init(3,4) in <o>+(x,y) % => 7
```

Since `init` returns `self`, the inner `let` above could be written as follows:

```
let
  o = new c
in
  let
    o = .<o>init(3,4)
  in
    <o>+(x,y)
```

or even, with just the outer `let` binding for `c`:

```
<.<new c>init(3,4)>+(x,y)
```

In the previous example, the `init` method is invoked separately, after the object is created, and not as part of the object creation itself. Naming this procedure “`init`” is not a requirement. A class can have several methods that initialize its fields, much as a Java class can have several constructors. Unlike Java, the OBJ language can apply its methods – even the ones intended to initialize the fields – at any time.

```
define c =  
  class  
    field x  
    method init = proc() {set x = 5 ; self}  
    method foo = proc() {set x = add1(x) ; self}  
  end
```

```
define o = .<new c>init()  
<o>x           % => 5  
<o>{.foo() ; x} % => 6  
<o>{.init() ; x } % => 5
```

The OBJ language has three additional expressions, with the following grammar rules:

```
<exp>:DisplayExp ::= DISPLAY <exp>  
                                     DisplayExp(Exp exp)  
<exp>:Display1Exp ::= DISPLAY1 <exp>  
                                     Display1Exp(Exp exp)  
<exp>:NewlineExp ::= NEWLINE  
                                     NewlineExp()
```

The DISPLAY, DISPLAY1, and NEWLINE tokens are defined by

```
DISPLAY 'display'  
DISPLAY1 'display#'  
NEWLINE 'newline'
```

Evaluating a DisplayExp expression results in evaluation of its <exp> part in the current environment; this value's toString() representation is then displayed on standard output, and the value is returned as the value of the DisplayExp expression. Display1Exp is like DisplayExp except that the displayed value is followed by a single space. The value of a NewlineExp expression is nil, and a newline is displayed on standard output.

Here are examples of how to use `display`, `display#`, and `newline`:

```
let
  x = 3
  y = 5
  z = 8
in
  { display x ; newline
    ; display y ; newline
    ; display z ; newline
  } % => returns nil
```

Evaluating this expression results in the following standard output (omitting the final `nil`):

```
3
5
8
```

If the `newline` expressions were not present, the output would appear as follows (omitting the final `8`):

```
358
```

If `display` were then replaced by `display#`, the output would appear as follows:

```
3 5 8
```

Consider the following OBJ program:

```
define summer =  
  class  
    field sum  
    method init = proc() {set sum = 0 ; self}  
    method add = proc(t) {set sum = +(sum,t) ; self}  
    method show = proc() {display sum ; newline ; self}  
  end
```

Here's an example of how this class might be used to find and display the sum of the integers 1, 3, 5, and 7:

```
define o = .<new summer>init()  
.<o>add(1)  
.<o>add(3)  
.<o>add(5)  
.<o>add(7)  
.<o>show()
```

Since `init` and `add` both return `self`, the following “one-liner” would accomplish the same thing:

```
.<.<.<.<.<.<.<new summer>init()>add(1)>add(3)>add(5)>add(7)>show()
```

Since we often encounter expressions like the following:

```
.<.<.<.<.<.<.<new summer>init()>add(1)>add(3)>add(5)>add(7)>show()
```

we introduce a short-hand way of writing this:

```
!<new summer>init()>add(1)>add(3)>add(5)>add(7)>show()!>
```

The token `LLANGLE`, defined as the string ‘!<’, introduces an expression – which must evaluate to an environment (class or object) – followed by a sequence of zero or more procedure applications. Each procedure is evaluated in the environment of the previous class or object, and the procedure application itself must return another object, whose environment is then used to evaluate the next procedure, and so on. The entire expression is terminated by a `RRANGLE` token, defined as ‘!>’. Note that the actual parameter expressions are evaluated in the environment in which the entire expression appears.

Here are the associated grammar rules:

`<exp>:EenvExp ::= LANGLE <exp> <mangle> RANGLE`

`EenvExp(Exp exp, Mangle mangle)`

`<mangle> **= RANGLE <exp> LPAREN <rands> RPAREN`

`Mangle(List<Exp> expList, List<Rands> randsList)`

The BNF identifier `mangle` should suggest “multiple angle (brackets)”, but it also could also appropriately be interpreted as a twisted mess.

```
<exp>:EenvExp      ::= LANGLE <exp> <mangle> RANGLE  
                        EenvExp(Exp exp, Mangle mangle)  
<mangle>:Mangle    **= RANGLE <exp> LPAREN <rands> RPAREN  
                        Mangle(List<Exp> expList, List<Rands> randsList)
```

Here is the implementation of how to evaluate an EenvExp expression.

```
EenvExp  
%%  
    public Val eval(Env env) {  
        Val v = exp.eval(env);      // the environment object  
        return mangle.eval(v, env);  
    }  
%%
```

The expression `exp` is evaluated in the current environment. Its value `v`, along with the current environment, is passed to the `mangle` object, which then evaluates the subsequent procedure applications.

`<exp>:EenvExp ::= LANGLE <exp> <mangle> RRANGLE`

`EenvExp(Exp exp, Mangle mangle)`

`<mangle>:Mangle **= RANGLE <exp> LPAREN <rands> RPAREN`

`Mangle(List<Exp> expList, List<Rands> randsList)`

A mangle object consists of a list of Exp objects and a list of Rands objects.

Here is the code for `eval (Val v, Env env)` in the `Mangle` class:

```
public Val eval(Val v, Env env) {
    Iterator<Exp> expIter = expList.iterator();
    Iterator<Rands> randsIter = randsList.iterator();
    while (expIter.hasNext()) {
        // expIter.next() ProcExp to apply
        // v.env() is the environment in which to build the ProcVal
        Val p = expIter.next().eval(v.env());
        // evaluate this method's rands in env
        List<Val> valList = randsIter.next().evalRands(env);
        v = p.apply(valList);
    }
    return v;
}
```

Each `Exp` is evaluated in the environment defined by the value `v` (which must, perforce, be a class or object), and this must evaluate to a `ProcVal` (named `p` in this code). The `valList` arguments to `p` are evaluated in the outside environment `env` (*not* in the environment defined by `v`) from the corresponding `Rands` object. The procedure `p` is then applied to these arguments, and the resulting application becomes the new `v`. This is repeated until all of the `Mangle` applications are performed. The final value `v` is returned as the result of the `EenvExp` expression.

It turns out that exposing the environment of a procedure can be used to implement, with procedures alone, a simplified approach to objects and methods. See the file

`/usr/local/pub/plcc/Code/OBJ/Prog/pobj`

for an example. On the other hand, exposing the environment of a procedure can result in modifying that environment, which can lead to unintended consequences.

In the OBJ language, the `send` operator provides syntactic sugar that is specific to method calls on an object. Specifically, for an object `o` and method `m`, the following expressions are equivalent:

```
send o m ( . . . )  
.<o>m ( . . . )
```

Some examples appear on the following slide.

Language OBJ (continued)

5.38

```
define c1 = class
  field x field y
  method init1 = proc(v,w) {set x=v ; set y=w}
  method getx1 = proc() x
  method gety1 = proc() y
end
```

```
define c2 = class extends c1
  field y
  method init2 = proc(w) set y=w
  method getx2 = proc() x
  method gety2 = proc() y
end
```

```
define o2 = new c2
```

```
send o2 init1(1,2) % same as .<o2>init1(1,2)
send o2 init2(9)   % same as .<o2>init2(9)
send o2 getx1()    % same as .<o2>getx1() => 1
send o2 gety1()    % same as .<o2>gety1() => 2
send o2 getx2()    % same as .<o2>getx2() => 1
send o2 gety2()    % same as .<o2>gety2() => 9
```

In method application, `self` always refers to the base object, even if `self` appears in the definition of a superclass method. This is what makes dynamic dispatch work!

```
define c1 =  
  class  
    method m1 = proc() 1  
    method m2 = proc() send self m1()  
  end  
define c2 =  
  class extends c1  
    method m1 = proc() 2  
  end  
define o1 = new c1  
define o2 = new c2  
  
send o1 m1() % same as .<o1>m1() => 1  
send o2 m1() % same as .<o2>m1() => 2  
send o2 m2() % same as .<o2>m2() => 2!
```

The following slide gives another example of how dynamic dispatch works.

Language OBJ (continued)

5.40

```
define shape =
  class
    method area = proc() 0 % shapeless
  end

define rectangle =
  class extends shape
    field l % length
    field w % width
    method init = proc(l,w) {set <this>l=l ; set <this>w=w ; self}
    method area = proc() *(l,w)
  end

define circle =
  class extends shape
    field rad % radius
    method init = proc(rad) {set <this>rad=rad ; self}
    method area = proc() *(3,*(rad,rad)) % a bit of an underestimate
  end

define r = .<new rectangle>init(4,5) % a rectangle with length 4 and width 5
define c = .<new circle>init(2)      % a circle with radius 2
define s = new shape

.<r>area() % => 20
.<c>area() % => 12
.<s>area() % => 0
```

Other examples on this slide and the next ...

```
define c1 =  
  class  
    method m1 = proc() send self m2()  
    method m2 = proc() 13  
  end  
define c2 =  
  class extends c1  
    method m1 = proc() 22  
    method m2 = proc() 23  
    method m3 = proc() send super m1()  
  end  
define c3 =  
  class extends c2  
    method m1 = proc() 32  
    method m2 = proc() 33  
  end  
define o3 = new c3  
  
send o3 m3()  % returns 33
```

Language OBJ (continued)

5.42

```
define a = class
  field i field j
  method setup = proc() {set i=15; set j=20; 50}
  method f = proc() .<self>g()
  method g = proc() +(i,j)
end
define b = class extends a
  field j field k
  method setup =
    proc() {set j=100; set k=200; .<super>setup(); .<self>h()}
  method g = proc() [i,j,k]
  method h = proc() .<super>g()
end
define c = class extends b
  method g = proc() .<super>h()
  method h = proc() +(j,k)
end
let
  p = proc(o)
    let
      u = .<o>setup()
    in
      [u, .<o>g(), .<o>f()]
in
  [.p(new a), .p(new b), .p(new c)]

% returns [[50,35,35],[35,[15,100,200],[15,100,200]],[300,35,35]]
```

In many object-oriented programming languages, the fields of an object can be made *private* – that is, inaccessible outside of the object's methods. For private fields, special publically accessible methods can be used to retrieve the field values or to modify them. These methods are often called *getters* and *setters*, respectively.

Suppose, for example, we provided a special designator called `private` that served to identify a field whose value was inaccessible outside of the class methods. Consider the following code:

```
define c =  
  class  
    private x  
    method get_x = proc() x  
    method set_x = proc($) set x = $  
  end  
define cc = new c  
.<cc>set_x(5) % ok - sets value of x to 5  
.<cc>get_x() % ok - returns 5  
<cc>x      % illegal - x is private
```

While this sort of code is common, there are two problems with this approach. The first is that every private field we want to access needs a getter and a setter. The second is that code such as `<cc>set x = 5` does not work, but is easier to write and understand than `.<cc>set_x(5)`.

The C# language championed by Microsoft solves these problems using the notion of a *property*. A property acts like a field but it provides built-in getter and setter code. When the field is *accessed*, the getter code is executed; when the field is *assigned to* with a `set` statement, the setter code is executed.

Here is the same class as described on the previous slide, with a property instead of a getter and setter:

```
define c =  
  class  
    field x  
    property x = prop x:set x=$  
  end  
define cc = new c  
<cc>set x = 5 % ok - sets the field value to 5  
<cc>x          % returns 5
```

The property `x` shadows the field `x`. This means that the field `x` cannot be accessed directly except through the property.

Here are the grammar rules for defining properties in a class definition:

```
<props>  **= PROPERTY <VAR> EQUALS <prop>  
          Props(List<Token> varList, List<Prop> propList)  
<prop>   ::= PROP <exp>getExp COLON <exp>setExp  
          Prop(Exp getExp, Exp setExp)
```

When a variable bound to a property is evaluated, its `getExp` code is executed using the environment captured where the property is defined, and the expressed value of the variable is the result of evaluating the `getExp` expression.

When a variable bound to a property is assigned to in a `set` expression, the RHS of the expression is evaluated in the current environment. The environment where the property is defined is then extended by binding the symbol `$` to the value of the RHS. The property's `setExp` expression is then evaluated in this extended environment, and the resulting value is the value of the `setExp` expression.

If a variable `z` [for example] is bound to a property whose `set` expression has no side-effects [such as `nil`], an expression such as `set z = ...` does not modify anything – including `z`.

```
define c =  
  class  
    field x  
    method init = proc(x) {set <this>x = x ; self}  
    property y = prop x:set x=$  
    property z = prop x:nil  
  end  
define o = .<new c>init(3)  
<o>[x,y,z]    % [3,3,3]  
<o>set x = 5  
<o>[x,y,z]    % [5,5,5]  
<o>set y = 11  
<o>[x,y,z]    % [11,11,11]  
<o>set z = 42  
<o>[x,y,z]    % [11,11,11]
```

We implement properties in the same way we implement call-by-name: a property evaluates to a thunk-like object called a `PropRef` (which extends the `Ref` class); this object captures the environment in which the property is defined, along with the property's `get` and `set` expressions. The expressed value of a variable bound to a property is the value obtained by evaluating the property's `get` expression in the captured environment by calling the thunk's `deRef` method. Similarly, when assigning a value to a variable bound to a property, the thunk's `setRef` method is called, with the assigned value bound to the special variable `$` and returning the value of the property's `set` expression.

Since the PROP language uses call-by-reference parameter passing semantics, the callee sees the `PropRef` instead of a copy of the variable's value, so any changes to the formal parameter bound to such a property directly affect the actual parameter.

```
define c =  
  class  
    static p = proc(t) set t = add1(t)  
    field x  
    method init = proc(x) {set <this>x = x ; self}  
    property x = prop x : set x = $  
    property y = prop x : set x = +($,$)  
  end  
define o = .<new c>init(3)  
<o>{.p(x) ; x} % evaluates to 4  
<o>{.p(y) ; x} % evaluates to 10
```

So far, a property is only defined in the context of a class definition, where it plays a role in object instantiation. It turns out that the behavior of properties could be useful even outside of the context of an object, especially to manage access to variables defined in a `let` expression. To make this explicit, we create a `letprop` construct that has the following concrete syntax and abstract class structure:

```
<exp>:LetpropExp ::= LETPROP <letpropDecls> IN <exp>  
                  LetpropExp(LetpropDecls letpropDecls, Exp exp)  
<letpropDecls>  **= <VAR> EQUALS <prop>  
                  LetpropDecls(List<Token> varList, List<Prop> propList)
```

Consider this (somewhat strange) example:

```
let  
  x = 3  
in  
  letprop  
    x = prop x : set x = 5  
  in  
    {set x = 42 ; x} % => 5
```

This expression evaluates to 5.

The `set` part of a `prop` is optional. If the `set` part is omitted, any attempt to apply the `set` operator to the variable results in a runtime exception. In this way, we can implement “read-only” properties.

```
let
  x = 3
in
  letprop
    x = prop x %% no 'set' part, so x is read only
  in
    {set x = 42 ; x} % => runtime exception
```

It turns out that a read-only `prop` behaves just like a `thunk` in call-by-name (which, you may recall, is also read-only), so we have the benefit of call-by-name semantics when we want it!

```
let
  while = proc(test?, do, ans)
    letrec
      loop = proc()
        if test? then {do ; .loop()} else ans
    in
      .loop()
  sum = 0
  count = 10
  i = 1
in
  letprop
    test? = prop count
    do = prop { set sum = +(sum, i)
               ; set i = add1(i)
               ; set count = sub1(count)
             }
    ans = prop sum
  in
    .while(test?, do, ans) % => 55 = 1+2+...+10
```